

Whole-Genome Regression and Prediction Methods Applied to Plant and Animal Breeding

Gustavo de los Campos,^{*,†} John M. Hickey,[†] Ricardo Pong-Wong,[‡] Hans D. Daetwyler,[§] and Mario P. L. Calus^{**}

^{*}Department of Biostatistics, School of Public Health, University of Alabama, Birmingham, Alabama 35294, [†]School of Environmental and Rural Science, University of New England, Armidale 2351, New South Wales, Australia, [‡]The Roslin Institute, Royal (Dick) School of Veterinary Studies, University of Edinburgh, Easter Bush, Midlothian, EH25 9RG, Scotland, [§]Biosciences Research Division, Department of Primary Industries, Bundoora 3083, Victoria, Australia, and ^{**}Animal Breeding and Genomics Centre, Wageningen UR Livestock Research, 8200 AB Lelystad, The Netherlands

ABSTRACT Genomic-enabled prediction is becoming increasingly important in animal and plant breeding and is also receiving attention in human genetics. Deriving accurate predictions of complex traits requires implementing whole-genome regression (WGR) models where phenotypes are regressed on thousands of markers concurrently. Methods exist that allow implementing these *large-p* with *small-n* regressions, and genome-enabled selection (GS) is being implemented in several plant and animal breeding programs. The list of available methods is long, and the relationships between them have not been fully addressed. In this article we provide an overview of available methods for implementing parametric WGR models, discuss selected topics that emerge in applications, and present a general discussion of lessons learned from simulation and empirical data analysis in the last decade.

MODERN animal and plant breeding schemes select individuals based on predictions of genetic merit. Rapid genetic progress requires that such predictions are accurate and that they can be produced early in life. Family-based predictions of genetic values have been used successfully for selection in plants and animals for many decades; however, there is a limit on the annual rate of genetic progress that can be attained with family-based prediction. Molecular markers allow describing the genome of individuals at a large number of loci, and this opens possibilities to derive accurate prediction of genetic values early in life. The first attempts to incorporate marker information into predictions were based on the presumption that one can localize causative mutations underlying genetic variation. This approach, known as QTL mapping (Soller and Plotkin-Hazan 1977; Soller 1978), led to the discovery of a few genes associated to genetic differences of traits of commercial interest. However, the impact on practical breeding programs has been smaller than initially envisaged

(Dekkers 2004; Bernardo 2008). Several factors contributed to this: first, with a few exceptions, the proportion of the variance accounted by mapped QTL has commonly been small. Second, the financial resources required to develop the populations needed to map QTL were considerable, limiting the adoption of this technology (see Dekkers 2004, Bernardo 2008, Collard and Mackill 2008, and Hospital 2009 for insightful discussions of lessons learned from QTL studies in animal and plant breeding).

There is a general consensus that most traits are affected by large numbers of small-effect genes (see Buckler *et al.* 2009, for examples of traits in maize affected by large numbers of small-effect loci) and that the prediction of complex traits requires considering large numbers of variants concurrently. The continued advancement of high-throughput genotyping and sequencing technologies allowed the discovery of hundreds of thousands of genetic markers (e.g., single-nucleotide polymorphisms, SNPs) in the genomes of humans and several plant and animal species. Such dense panels of molecular markers allow exploiting multilocus linkage disequilibrium (LD) between QTL and genome-wide markers (e.g., SNPs) to predict genetic values. Although earlier contributions exist (Nejati-Javaremi *et al.* 1997; Haley and Visscher 1998; Whittaker *et al.* 2000), the foundations of genome-enabled selection (GS) were largely defined in the ground-breaking article by Meuwissen *et al.* (2001),

Copyright © 2013 by the Genetics Society of America

doi: 10.1534/genetics.112.143313

Manuscript received May 18, 2012; accepted for publication June 11, 2012

Available freely online through the author-supported open access option.

Supporting information is available online at <http://www.genetics.org/content/suppl/2012/06/28/genetics.112.143313.DC1>.

[†]Corresponding author: University of Alabama, 1665 University Blvd., 327L Ryals Public Health Bldg., Birmingham, AL 35216. E-mail: gcampos@uab.edu

who proposed to incorporate dense molecular markers into models using a simple, but powerful idea: regress phenotypes on all available markers using a linear model. And in recent years this approach has gained ground both in animal (VanRaden *et al.* 2009) and plant breeding (Bernardo and Yu 2007; Crossa *et al.* 2010).

With high-density SNP panels the number of markers (p) can vastly exceed the number of records (n), and fitting this *large-p with-small-n* regression requires using some type of variable selection or shrinkage estimation procedure. Owing to developments of penalized and Bayesian estimation procedures, as well as advances in the field of nonparametric regressions, several shrinkage estimation methods have been proposed and used for whole-genome regression (WGR) and prediction (WGP) of phenotypes or breeding values. However, the relationships between these methods have not been fully addressed and many important topics emerging in empirical applications have been often overlooked. In this article we provide an overview of parametric Bayesian methods as applied to GS (*Methods*), a discussion of selected topics that emerge when these models are used for empirical analysis (*Selected Topics Emerging in Empirical Applications*), and a discussion of lessons learned in the past years based on a literature review of simulation and empirical studies (*Lessons Learned from Simulation and Empirical Data Analysis*).

Methods

Early proposals for implementing GS (Meuwissen *et al.* 2001) used linear regression methods. More generally, one can regress phenotypes on marker covariates using a regression function, $f(x_{i1}, x_{i2}, \dots, x_{ip})$ that may be parametric or not so that $y_i = f(x_{i1}, x_{i2}, \dots, x_{ip}) + \varepsilon_i$. Here, the regression function, $f(x_{i1}, x_{i2}, \dots, x_{ip})$, should be viewed as an approximation to the true unknown genetic values, $\{g_i\}_{i=1}^n$, which can be a complex function involving the genotype of the i th individual at a large number of genes as well as cryptic interactions between genes and between genes and environmental conditions. Therefore, in a WGR model residuals $\{\varepsilon_i\}_{i=1}^n$ represent random variables capturing nongenetic effects, plus approximation errors, $g_i - f(x_{i1}, x_{i2}, \dots, x_{ip})$, which can emerge due to imperfect LD between markers and QTL or because of model misspecification (e.g., unaccounted interactions).

The *linear model* appears as a special case with $f(x_{i1}, x_{i2}, \dots, x_{ip}) = \mu + \sum_{j=1}^p x_{ij}\beta_j$, where μ is an intercept, x_{ij} is the genotype of the i th individual at the j th marker ($j = 1, \dots, p$), and β_j is the corresponding marker effect. Alternatively the regression function could be represented using semiparametric approaches (Gianola *et al.* 2006; de los Campos *et al.* 2010a) such as reproducing kernel Hilbert spaces (RKHS) regressions or neural networks (NN). Therefore, a first element of model specification in WGR is whether genetic values are approximated by using linear regression procedures or using semiparametric methods. In this article we focus on linear regression models; a review

and a discussion about nonparametric procedures can be found in Gianola *et al.* (2010).

With modern genotyping technologies the number of markers, and therefore the number of parameters to be estimated, can vastly exceed the number of records. To confront the problems emerging in these large- p with small- n regressions, estimation procedures performing variable selection, shrinkage of estimates, or a combination of both are commonly used. Therefore, a second element of model choice pertains to the type of shrinkage estimation procedure used. Next, we discuss briefly the effects of shrinkage on the statistical properties of estimates and subsequently review some of the most commonly used penalized and Bayesian variable selection and shrinkage estimation procedures.

Effects of shrinkage on the mean-squared error of estimates

The accuracy of an estimator can be measured with the squared Euclidean distance between the estimated, $\hat{\theta}$, and the true value of the parameter, θ . In the case of scalars this is simply the squared deviation: $\|\hat{\theta}(\mathbf{y}) - \theta\|^2 = [\hat{\theta}(\mathbf{y}) - \theta]^2$. Here, we write $\hat{\theta}(\mathbf{y})$ to stress that the estimator is a function of the sampled data. The mean-squared error (MSE) is the expected value (over possible realizations of the data) of the squared Euclidean distance, $\text{MSE}(\hat{\theta}) = E[\hat{\theta}(\mathbf{y}) - \theta]^2$. This can be decomposed into two terms: the variance of the estimator plus the square of its bias, $\text{MSE}(\hat{\theta}) = \text{Var}[\hat{\theta}] + \text{Bias}[\hat{\theta}]^2 = E\{[\hat{\theta} - E(\hat{\theta})]^2\} + [E(\hat{\theta}) - \theta]^2$. The variance of the estimator (and in some cases its bias) decreases with sample size. With standard estimation procedures, such as ordinary least squares (OLS) or maximum likelihood (ML), with fixed sample size, the variance of estimates increases rapidly as p does, yielding high MSE of estimates. One way of confronting the MSE problem emerging in large- p with small- n regressions is by shrinking estimates toward a fixed point (e.g., 0); this may increase bias but reduces the variance of the estimator. To illustrate the effects of shrinkage on MSE of estimates, consider a simple shrinkage estimator obtained by multiplying an unbiased estimator $\hat{\theta}$ by a constant $\alpha \in [0, 1]$ so that $\tilde{\theta} = \alpha\hat{\theta} + (1-\alpha)0 = \alpha\hat{\theta}$. The new estimator shrinks the original one toward 0. If $\theta \neq 0$, $\tilde{\theta}$ is biased; however, the variance of the new estimator, $\text{Var}(\alpha\hat{\theta}) = \alpha^2\text{Var}(\hat{\theta})$ is guaranteed to be lower for any $\alpha < 1$. Penalized and Bayesian methods are the two most commonly used shrinkage estimation procedures, an overview of these methods is given next.

Penalized methods

In penalized regressions estimates are derived as the solution to an optimization problem that balances model goodness of fit to the training data and model complexity. For continuous outcomes, lack of fit to the training data is usually measured by the residual sum of squares, $\sum_i (y_i - \mu - \sum_{j=1}^p x_{ij}\beta_j)^2$ (alternatively, one can use the negative of the logarithm of the likelihood or some other loss function) and model complexity is commonly defined as

a function of model unknowns, $J(\boldsymbol{\beta})$; therefore, penalized estimates are commonly derived as the solution to an optimization problem of the form

$$(\hat{\mu}, \hat{\boldsymbol{\beta}})_{\text{argmin}} \left\{ \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 + \lambda J(\boldsymbol{\beta}) \right\}, \quad (1)$$

where $\lambda \geq 0$ is a regularization parameter that controls the trade-offs between lack of fit and model complexity. Ordinary least squares appear as a special case of (1) with $\lambda = 0$. Usually, not all model unknowns are penalized; for instance, in (1) the intercept is not included in the penalty function. The features of the regression function that are not penalized [the overall mean in (1)] are then perfectly fitted.

Several penalized estimation procedures have been proposed, and they differ on the choice of penalty function, $J(\boldsymbol{\beta})$. In ridge regression (RR) (Hoerl and Kennard 1970), the penalty is proportional to the sum of squares of the regression coefficients or L2 norm, $J(\boldsymbol{\beta}) = \sum_{j=1}^p \beta_j^2$. A more general formulation, known as bridge regression (Frank and Friedman 1993), uses $J(\boldsymbol{\beta}) = \sum_{j=1}^p \|\beta_j\|^\gamma$ with $\gamma > 0$. RR is a particular case with $\gamma = 2$ yielding the L2 norm, $J(\boldsymbol{\beta}) = \sum_{j=1}^p \|\beta_j\|^2$. Subset selection occurs as a limiting case with $\gamma \rightarrow 0$, which penalizes the number of nonzero effects regardless of their magnitude, $J(\boldsymbol{\beta}) = \sum_{j=1}^p 1(\beta_j \neq 0)$. Another special case, known as least absolute angle and selection operator (LASSO) (Tibshirani 1996), occurs with $\gamma = 1$, yielding the L1 penalty: $J(\boldsymbol{\beta}) = \sum_{j=1}^p \|\beta_j\|$. Using this penalty induces a solution that may involve zeroing out some regression coefficients and shrinkage estimates of the remaining effects; therefore LASSO combines variable selection and shrinkage of estimates. LASSO has become very popular in several fields of applications. However, LASSO and subset selection approaches have two important limitations. First, by construction, in these methods the solution admits at most n nonzero estimates of regression coefficients (Park and Casella 2008). In WGR of complex traits, there is no reason to restrict the number of markers with nonzero effect to be limited by n (the number of observations). Second, when predictors are correlated, something that occurs when LD span over long regions, methods performing variable selection such as the LASSO are usually outperformed by RR (Hastie *et al.* 2009). Therefore, in an attempt to combine the good features of RR and of LASSO in a single estimation framework, Zou and Hastie (2005) proposed to use as a penalty a weighted average of the L1 and L2 norm, that is, for $0 \leq \alpha \leq 1$, $J(\boldsymbol{\beta}) = \alpha \sum_{j=1}^p \|\beta_j\| + (1-\alpha) \sum_{j=1}^p \|\beta_j\|^2$, and termed the method the elastic net (EN). This model involves two tuning parameters that need to be specified, the regularization parameter (λ) and α .

Bayesian shrinkage estimation

Bayesian methods can also be used for variable selection and shrinkage of estimates. Most penalized estimates are equivalent to posterior modes of certain types of Bayesian models (Kimeldorf and Wahba 1970; Tibshirani 1996). An illustration

of the equivalence between some penalized and some Bayesian estimates is given in supporting information, File S1.

The general structure of the standard Bayesian linear models used in GS is

$$\begin{aligned} p(\mu, \boldsymbol{\beta}, \sigma^2 | \mathbf{y}, \boldsymbol{\omega}) & \propto p(\mathbf{y} | \mu, \boldsymbol{\beta}, \sigma^2) p(\mu, \boldsymbol{\beta}, \sigma^2 | \boldsymbol{\omega}) \\ & \propto \prod_{i=1}^n N(y_i | \mu + \sum_{j=1}^p x_{ij} \beta_j, \sigma^2) \prod_{j=1}^p p(\beta_j | \boldsymbol{\omega}) p(\sigma^2), \end{aligned} \quad (2)$$

where $p(\mu, \boldsymbol{\beta}, \sigma^2 | \mathbf{y}, \boldsymbol{\omega})$ is the posterior density of model unknowns $\{\mu, \boldsymbol{\beta}, \sigma^2\}$ given the data (\mathbf{y}) and hyperparameters ($\boldsymbol{\omega}$), $p(\mathbf{y} | \mu, \boldsymbol{\beta}, \sigma^2) = \prod_{i=1}^n N(y_i | \mu + \sum_{j=1}^p x_{ij} \beta_j, \sigma^2)$ is the conditional density of the data given the unknowns, which for continuous traits are commonly independent normal densities with mean $E(y_i | \mu, \boldsymbol{\beta}, \sigma^2) = \mu + \sum_{j=1}^p x_{ij} \beta_j$ and with variance $\text{Var}(y_i | \mu, \boldsymbol{\beta}, \sigma^2) = \sigma^2$, and $p(\mu, \boldsymbol{\beta}, \sigma^2 | \boldsymbol{\omega}) \propto \prod_{j=1}^p p(\beta_j | \boldsymbol{\omega}) p(\sigma^2)$ is the joint prior density of model unknowns, including the intercept (μ), which is commonly assigned a flat prior, marker effects $\{\beta_j\}$, which are commonly assigned IID informative priors, and the residual variance (σ^2), which is commonly assigned a scaled-inverse chi-square prior with degree of freedom d.f. and scale parameter S , which is $p(\sigma^2) = \chi^{-2}(\sigma^2 | \text{d.f.}, S)$; here we use a parameterization $E(\sigma^2 | \text{d.f.}, S) = (\text{d.f.} \times S) / (\text{d.f.} - 2)$.

In these Bayesian models, the prior density of marker effects, $p(\beta_j | \boldsymbol{\omega})$, defines whether the model will induce variable selection and shrinkage or shrinkage only. Also, the choice of prior will define the extent and type of shrinkage induced. Two important features of these priors are how much mass they have in the neighborhood of zero and how thick or flat the tails of the density are. Based on these two features we classified the most commonly used priors into four big categories and in Figure 1 we have arranged them in a way that, starting from the Gaussian prior located in the top left corner, as one moves clockwise there is an increase in the peak of mass at zero and the tails are allowed to become thicker.

Gaussian prior: This density (depicted in the top left corner of Figure 1) has two hyperparameters: the mean (commonly set to zero) and the variance (σ_β^2); therefore, in this model, $\boldsymbol{\omega} = \sigma_\beta^2$. If the intercept and the variance parameters are known, the posterior distribution of marker effects, $p(\boldsymbol{\beta} | \mathbf{y}, \mu, \sigma^2, \sigma_\beta^2) \propto \prod_{i=1}^n N(y_i | \mu + \sum_{j=1}^p x_{ij} \beta_j, \sigma^2) \prod_{j=1}^p N(\beta_j | 0, \sigma_\beta^2)$, can be shown to be multivariate normal, with posterior mean given by $\hat{\boldsymbol{\beta}} = [\mathbf{X}'\mathbf{X} + \sigma^2 \sigma_\beta^{-2} \mathbf{I}]^{-1} \mathbf{X}'\tilde{\mathbf{y}}$, where $\mathbf{X} = \{x_{ij}\}$ is a matrix of marker genotypes and $\tilde{\mathbf{y}} = \{y_i - \mu\}$ is a vector of (centered) phenotypes. This is exactly the RR estimate with $\lambda = \sigma^2 / \sigma_\beta^2$. Because of this characteristic, this model is sometimes referred to as Bayesian ridge regression (BRR). Also, $\hat{\boldsymbol{\beta}}$ can be shown to be the best linear unbiased predictor (BLUP) of marker effects; therefore, this model is also sometimes referred to as RR-BLUP (standing for ridge regression BLUP).

Ridge regression and the BRR both perform an extent of shrinkage that is homogenous across markers; this approach




Figure 1 Commonly used prior densities of marker effects (all with zero mean and unit variance). The densities are organized in a way that, starting from the Gaussian in the top left corner, as one moves clockwise, the amount of mass at zero increases and tails become thicker and flatter.

may not be optimal if some markers are linked to QTL while others are in regions that do not harbor QTL. To overcome this potential limitation, other prior densities can be used.

Thick-tailed priors: The two most commonly used thick-tailed densities, the scaled t and the double exponential, are represented in the top right corner of Figure 1. The scaled- t density is the prior used in model BayesA (Meuwissen *et al.* 2001) and the double-exponential or Laplace prior is the density used in the Bayesian LASSO (BL) (Park and Casella 2008). Relative to the Gaussian, these densities have higher mass at zero (inducing strong shrinkage toward zero of estimates of effects of markers with small effects) and thicker tails (inducing, relative to the BRR, less shrinkage of estimates of markers with sizable effects).

For computational convenience, the thick-tail densities are commonly represented as infinite mixtures of scaled-normal densities (Andrews and Mallows 1974) of the form $p(\beta_j | \omega) = \int N(\beta_j | 0, \sigma_{\beta_j}^2) p(\sigma_{\beta_j}^2 | \omega) d\sigma_{\beta_j}^2$, where $p(\sigma_{\beta_j}^2 | \omega)$ is a prior density assigned to marker-specific variance parameters and ω denotes hyperparameters indexing this density.

When $p(\sigma_{\beta_j}^2 | \omega)$ is a scaled inverse chi-square density, the resulting marginal prior of marker effects is scaled t , and this is the approach used in model BayesA of Meuwissen *et al.* (2001). When $p(\sigma_{\beta_j}^2 | \omega)$ is an exponential density, the resulting marginal prior of marker effects is double exponential, and this is the approach followed in the Bayesian LASSO of Park and Casella (2008). The double-exponential density is indexed by a single parameter (rate) and the scaled t is indexed by two parameters (scale and degree of freedom); this gives the scaled t more flexibility for controlling how thick the tails may be. An even higher degree of flexibility to control the shape of the prior can be obtained by using priors that are finite mixtures.

Spike-slab priors: These models use priors that are mixtures of two densities: one with small variance (the spike) and one with large variance (the slab) (e.g., George and McCulloch 1993; see bottom left corner in Figure 1). Commonly, the spike and the slab are both zero-mean normal densities. A graphical representation of one of such mixtures is given in the bottom right corner of Figure 1.

The general form of these mixtures is $p(\beta_j | \pi, \sigma_{\beta_1}^2, \sigma_{\beta_2}^2) = \pi \times N(\beta_j | 0, \sigma_{\beta_1}^2) + (1 - \pi) \times N(\beta_j | 0, \sigma_{\beta_2}^2)$, where $\pi \in [0, 1]$ is a mixture proportion and $\sigma_{\beta_1}^2$ and $\sigma_{\beta_2}^2$ are variance parameters. To prevent the so-called label-switching problem a common approach is to restrict $\sigma_{\beta_1}^2 \leq \sigma_{\beta_2}^2$ so that π can be interpreted as the proportion of effects coming from the “small” variance component. Another approach is to reparameterize the prior so that the variance of one of the components is a scaled version of the variance of the other component; for instance, $p(\sigma_{\beta_j}^2 | \pi, \sigma_{\beta}^2, \tau) = \pi \times N(\beta_j | 0, \tau^{-1} \sigma_{\beta}^2) + (1 - \pi) \times N(\beta_j | 0, \sigma_{\beta}^2)$ with $\tau > 1$. Model Bayes “stochastic search variable selection” (SSVS) (Calus *et al.* 2008; Verbyla *et al.* 2009) follows this approach with τ commonly fixed at a certain value (e.g., $\tau = 100$). Another possibility is to link the proportion of effects coming from the “small-variance” component with the proportion of variance accounted for. For instance, following Yu and Meuwissen (2011), one could assume that $\tilde{\pi}\%$ of the markers account for $(100 - \tilde{\pi})\%$ of genetic variance. The so-called Pareto principle represents a specific case of the more general principle with $\tilde{\pi} = 20$. This reduces the number of hyperparameters that need to be chosen, at the expense of imposing restrictions that may or may not hold.

Although spike-slab models are usually formed by mixing two Gaussian components, similar models may be obtained by mixing other densities such as scaled t (Zou *et al.* 2010) or double exponential (DE). There are two limiting cases of the spike-slab model that are of special interest. The first one occurs when $\pi = 0$ or $\pi = 1$, and this case corresponds to the standard Gaussian prior (see above); the second one occurs when $\sigma_{\beta_1}^2 \rightarrow 0$, and in this case, the small-variance component of the mixture collapses to a point of mass at zero, giving rise to a prior that consists of a mixture of a point of mass at zero and a slab.

Point of mass at zero and slab priors: These are used to induce a combination of variable selection and shrinkage (see bottom left corner in Figure 1). These priors are used, for example, in models BayesB (Meuwissen *et al.* 2001) and BayesC (Habier *et al.* 2011). In Bayes B the slab is a scaled- t density, while in BayesC the slab is a normal density.

Genome-enabled BLUP

An alternative parameterization of the BRR can be obtained by replacing $\sum_{j=1}^p x_{ij} \beta_j$ with $u_i = \sum_{j=1}^p x_{ij} \beta_j$ or in matrix notation $\mathbf{u} = \mathbf{X}\boldsymbol{\beta}$. In the BRR, marker effects are IID normal random variables. From properties of the multivariate normal density it follows that $\mathbf{u} \sim N(\mathbf{0}, \mathbf{X}\mathbf{X}'\sigma_{\beta}^2) = N(\mathbf{0}, \mathbf{G}\sigma_u^2)$, where $\mathbf{G} = \mathbf{X}\mathbf{X}'k$ for some k . For instance, a common choice is to use $k^{-1} = 2\sum_{j=1}^p \theta_j(1 - \theta_j)$, where θ_j is (an estimate of) the frequency of the allele coded as one at the j th marker. Indeed, $\mathbf{G} = \mathbf{X}\mathbf{X}'k$ can be regarded as an estimate of the realized matrix of additive relationships (Habier *et al.* 2007; VanRaden 2007). Therefore, an equivalent represen-

tation of the BRR is given by the following model [genome-enabled BLUP, G-BLUP]:

$$p(\mathbf{u}, \mathbf{y} | \mu, \sigma^2, \sigma_{\beta}^2) \propto \prod_{i=1}^n N(y_i | \mu + u_i, \sigma^2) N(\mathbf{u} | \mathbf{0}, \mathbf{G}\sigma_u^2). \quad (3)$$

The posterior mode of this model can be shown to be

$$\hat{\mathbf{u}} = [\mathbf{I} + \lambda \mathbf{G}^{-1}]^{-1} \tilde{\mathbf{y}} = \mathbf{X}\hat{\boldsymbol{\beta}} \quad (4)$$

with $\lambda = \sigma^2 \sigma_u^{-2}$. This is also the best linear unbiased predictor of \mathbf{u} , and therefore this model is usually referred as to G-BLUP. Above, we have motivated G-BLUP by exploiting its equivalence with the BRR; however, these methods can be also motivated simply as an additive infinitesimal model in which we replace the standard pedigree-based numerator relationship matrix with a marker-based estimate of additive relationships. Indeed, these methods have existed long before GS emerged (Ritland 1996, 2002; Nejati-Javaremi *et al.* 1997; Lynch and Ritland 1999; Eding and Meuwissen 2001).

Computing genomic relationships for G-BLUP: Several proposals exist as to how to map from pairs of marker genotypes onto estimates of genetic relationships, and no one is considered superior. A first distinction is between methods that aim at estimating realized genomic relationships [or proportion of identical by state (IBS) (VanRaden 2007; Yang *et al.* 2010)] and those that attempt to estimate probability of sharing alleles due to inheritance from a known common ancestor [or probability of identical by descent (IBD)]. The IBD methods (Pong-Wong *et al.* 2001; Villanueva *et al.* 2005) are essentially multilocus extensions of the single-locus IBD approach of Fernando and Grossman (1989) with IBD coefficients averaged across multiple putative QTL.

Within the IBS framework, the most common approach, at least in GS, is to estimate genomic relationships using moment-based estimators, which in general take the form of cross-products of marker genotypes: $G_{ii} \propto \sum_{k=1}^p x_{ik} x_{ik}$. Here p is the number of loci and $x_{ij}, x_{i'j}$ are the genotypes of individuals i and i' at the j th locus. In matrix notation we have $\mathbf{G} \propto \mathbf{X}\mathbf{X}'$. As with any other regression procedure, marker genotypes can be centered by subtracting the mean of the marker genotype or centered and standardized to a unit variance; that is, $\tilde{x}_{ij} = (x_{ij} - 2\theta_j) / \sqrt{2\theta_j(1 - \theta_j)}$, where $\{\theta_j\}_{j=1}^p$ are estimates of the frequency of the allele coded as one. Therefore, another common estimator is $\tilde{G}_{ii} \propto \sum_{k=1}^p \tilde{x}_{ik} \tilde{x}_{ik}$. Centering implies that variances and covariances between genetic values are measured as deviations with respect to a center defined by the average genotype. Following the tradition of pedigree-based infinitesimal models, one can define the “center” to be the average genotype in a “base” population. In such case allele frequencies should be estimated in that population. Alternatively, allele frequencies could be estimated directly from the sample without further

consideration about an ancestral base population of nominally unrelated individuals. In this case the “origin” is defined as the average genotype in the sample. When this approach is used, some entries of \mathbf{G} may become negative, some diagonal elements become <1 , and the average diagonal value has an expected value equal to 1. Therefore, we cannot interpret the entries of \mathbf{G} as proportion of allele sharing or as probabilities. Nevertheless, from the point of view of the Gaussian process $G_{ii'}$ simply defines a covariance function and nothing precludes assigning negative prior covariances between pairs of genetic values. For \mathbf{G} to define a proper Gaussian process, it must be positive semidefinite; this is guaranteed when $\mathbf{G} \propto \mathbf{X}\mathbf{X}'$ or $\mathbf{G} \propto \mathbf{X}\mathbf{X}'$. However, other methods do not guarantee that this condition will hold. Therefore, a good practice is to check that this condition is satisfied, by, for example, checking that the associated eigenvalues of \mathbf{G} are all nonnegatives.

Relative to estimates of genomic relationship based on unstandardized markers, $G_{ii'}$, standardization, $\tilde{G}_{ii'}$, increases the “weight” given to markers with extreme allele frequency on the computation of genomic relationships and this occurs because the denominator used in the standardization, $\sqrt{2\theta_j(1-\theta_j)}$, is maximum at intermediate allele frequencies and minimum at extreme allele frequencies. Yang *et al.* (2010) proposed a modified version of $\tilde{G}_{ii'}$ where a different formula is used to compute the diagonal elements of \mathbf{G} . The proposed formula has a sampling variability that does not depend on allele frequency and equals one in absence of inbreeding; however, to the best of our knowledge, the proposed method is not guaranteed to yield a positive semidefinite matrix.

Relationships between Bayesian methods

The categorization of priors given in Figure 1 is somehow arbitrary and some models can be considered special or limiting cases of others. In Figure 2 we represent some of these relationships. (Figure 2 is partially inspired by a presentation given by Robert Tempelman at University of Alabama at Birmingham, who discussed connections between G-BLUP, BayesA, and BayesB.):

1. In finite mixture models we mix K densities; therefore, models using two or a single density component can be seen as special cases of the finite mixture model with $K = 2$ and $K = 1$, respectively (see paths 1a–1c in Figure 2).
2. Starting with a two-component mixture, such as the spike–slab, we can obtain models with a point of mass at zero and a slab, such as BayesB or –C, by fixing the variance of one of the components at zero (see paths 2a and 2b in Figure 2).
3. Models BayesA and BRR can be obtained as special cases of models BayesB and BayesC, respectively. This is done by setting in either BayesB or BayesC the proportion of markers with no effect (π) equal to zero (see paths 3a and 3b in Figure 2).
4. The scaled- t density has two parameters, the scale and the d.f.; as d.f. increases, the scaled- t density becomes

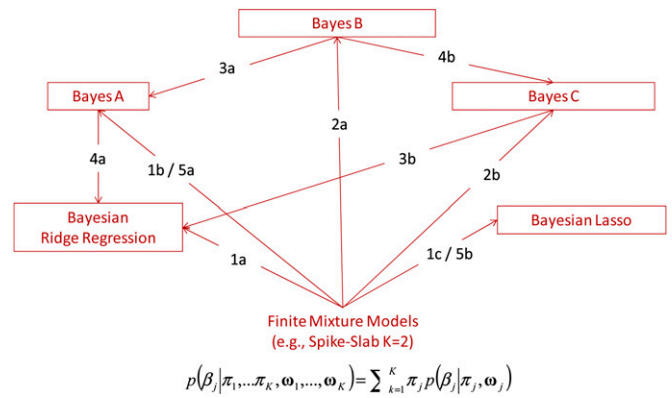


Figure 2 Relationships between some prior densities commonly assigned to marker effects.

increasingly similar to the Gaussian density. Therefore, starting from BayesA (BayesB) one can obtain the BRR (BayesC) by simply setting the d.f. to a very large value (see 4a and 4b in Figure 2).

5. The spike–slab prior is commonly formed by mixing two normal densities. The flexibility of such mixtures can be increased by increasing the number of components; eventually, we could consider an infinite number of components, each of which will have its own variance. However, there is a limit on the number of variance parameters that we can estimate. To confront this, a common approach is to regard these variances as random variables that are drawn from a common process. This is precisely what the scaled- t or DE densities are: these are infinite mixtures of scaled-normal densities (see paths 5a and 5b in Figure 2).

In view of the fact that many models (BRR, BayesA, and BayesC) appear as special cases of BayesB (for some values of parameters π and d.f.), a reasonable strategy would be to use a modified version of BayesB with π , scale and d.f. estimated from data (Nadaf *et al.* 2012). However, usually, with long-range LD and with $p \gg n$, different configurations of marker effects can yield very similar values at the likelihood. Therefore, estimating π , the scale and d.f. parameters jointly from the data may not be possible.

Dealing with hyperparameters

The hyperparameters indexing the prior density of marker effects (ω) control the extent and strength of shrinkage of estimates of marker effects and they can have important impacts on inferences; therefore, dealing with these hyperparameters appropriately is crucial. These unknowns can be dealt with in different ways.

Heritability-based rules: One possibility is to choose these hyperparameters based on prior expectation about the genetic variance of the trait. This approach was used, for instance, by Meuwissen *et al.* (2001), who derived hyperparameter values of the models by solving for them

as a function of genetic variance. In their derivation they assume that genetic variance emerges due to uncertainty both about genotypes and about marker effects; however, this is not entirely consistent with the Bayesian models used in GS where genotypes are regarded as fixed and marker effects as random. Therefore, here we present a simple derivation that is consistent with models used in GS (Equation 2) where marker genotypes are regarded as fixed and marker effects are viewed as random IID variables. In linear models we have $g_i = \sum_{j=1}^p x_{ij}\beta_j$; therefore the prior variance of the i th genomic value is $\text{Var}(g_i) = \text{Var}(\sum_{j=1}^p x_{ij}\beta_j) = [\sum_{j=1}^p x_{ij}^2] \times \text{Var}(\beta_j | \omega)$, where $\text{Var}(\beta_j | \omega)$ is the prior variance of marker effects that is a function of ω . Therefore, the average prior variance of genetic values in the sample is $V_g = [n^{-1} \sum_{i=1}^n \sum_{j=1}^p x_{ij}^2] \times \text{Var}(\beta_j | \omega) = \text{MS}_X \text{Var}(\beta_j | \omega)$, where $\text{MS}_X = n^{-1} \sum_{i=1}^n \sum_{j=1}^p x_{ij}^2$ is the average sum of squares of marker genotypes. Commonly, the model includes an intercept and variance is defined as deviations of genomic values from the center of the sample; therefore, MS_X should be computed using centered genotypes; that is, $\text{MS}_X = n^{-1} \sum_{i=1}^n \sum_{j=1}^p (x_{ij} - \bar{x}_j)^2$, where \bar{x}_j is the average genotype at the j th marker. Moreover, when marker genotypes are centered and standardized to a null mean and unit variance MS_X equals the number of markers (p). A natural approach is to replace V_g with the product of (an estimate of) heritability and of the variance of phenotypes, yielding

$$\text{Var}(\beta_j | \omega) = \frac{h^2 \sigma_p^2}{\text{MS}_X}. \quad (5)$$

Equation 5 can be used to solve for values of ω . We have only one equation; therefore, if ω involves more than one hyperparameter, others need to be fixed. Table 1 shows examples of the use of this formula for BRR, BayesA, BayesB, BayesC, Bayesian LASSO, and Bayes SSVS.

Validation methods: Another possibility is to fit models over a grid of values of ω and then retain the value that maximizes predictive performance. To that end some type of internal validation (e.g., using a tuning data set) needs to be carried out. However, this approach can be computationally demanding. This happens because the grid of values of ω may involve a large number of cells (especially when ω involves several parameters) and because standard validation schemes usually involve fitting the model several times (e.g., across different folds of a cross-validation) for each possible value of ω in the grid. Other alternatives such as leave-one-out cross-validation or generalized cross-validation could be used; however, unlike ordinary least squares, in many of the models of interest, the leave-one-out residual sum of squares does not have a closed form. Because of these reasons this approach has not been a very popular one in GS [although examples of the use of validation methods for choosing regularization parameters exist (Usai *et al.*, 2009)].

Fully Bayesian treatment: The fully Bayesian approach regards ω as unknown. This is done by assigning a prior to ω ; the model in expression (2) becomes

$$p(\mu, \beta, \sigma^2, \omega | y) \propto \prod_{i=1}^n N(y_i | \mu + \sum_{j=1}^p x_{ij}\beta_j, \sigma^2) \left\{ \prod_{j=1}^p p(\beta_j | \omega) p(\sigma^2) \right\} \times p(\omega | H), \quad (6)$$

where $p(\omega | H)$ is a prior density assigned to ω and H denotes a set of hyperparameters of higher order. The above-mentioned heritability rule can be seen as a limiting case of (6), where $p(\omega | H)$ is set to be simply a point of mass at some value, say ω_0 , which was chosen using prior knowledge (e.g., using the formulas in Table 1). In the fully Bayesian treatment, we may choose H so that $p(\omega | H)$ has a prior mean or prior mode in the neighborhood of ω_0 . This incorporates prior information into the model but in a more flexible way than heritability-based rules. The choice of prior, $p(\omega | H)$ depends on the nature of the hyperparameters. For BRR $\omega = \sigma_\beta^2$; therefore, it is natural to choose $p(\sigma_\beta^2 | H) = \chi^{-2}(\sigma_\beta^2 | \text{d.f.}_\beta, S_\beta)$; for BL Park and Casella (2008) suggested using a Gamma prior for λ^2 , $p(\lambda^2 | H) = G(\lambda^2 | \text{shape}, \text{rate})$; therefore, $p(\lambda | H) = G(\lambda^2 | \text{shape}, \text{rate}) 2\lambda$. For model BayesA $\omega = \{\text{d.f.}_\beta, S_\beta\}$, a common choice is to assign a Gamma prior to the scale parameter and either fix the degree of freedom parameter (usually to some small value >4 to guarantee a finite prior variance) or assign a prior to d.f._β with support on the positive real numbers.

Empirical Bayes methods: The empirical-based approach involves replacing $p(\omega | H)$ in (3) with a point of mass located at an estimate of ω ; that is, $p(\omega | H) = 1\{\omega = \hat{\omega}\}$. In this respect this approach is similar to heritability-based rules; however, in the empirical Bayes method (EB) $\hat{\omega}$ is a data-derived estimate. This approach is also commonly used in pedigree-based models where first, variance components are estimated from the data using restricted maximum likelihood and then, BLUPs of breeding values are derived with variance parameters replaced with those estimates. Ideally, we want $\hat{\omega}$ to be the posterior mean of ω ; however, in most cases it is difficult to derive a closed-form formula for the marginal posterior mean of ω . An alternative is to use the empirical Bayes principle within a Gibbs sampler (Casella 2001); however, the convergence of the algorithm may be too slow, and evidence does not suggest superiority of this approach relative to the fully Bayesian treatment.

Relaxing IID assumptions

All of the above-mentioned Bayesian models use IID priors for marker effects; that is, $p(\beta_j | \omega)$ is the same for all markers; therefore, the prior mean and variances are the same for all marker effects. This assumption can be justified based on “ignorance”; however, in many instances we may have additional prior information about markers and we may want to incorporate such information into the prior

Table 1 Prior density of marker effects, prior variance of marker effects, and suggested formulas for choosing hyperparameter values by model

Model $p(\beta_j \omega)$	Hyperparameters	Prior variance $\text{Var}(\beta_j \omega)$	Solution for scale/variance parameter
Bayesian ridge regression			
$N(\beta_j 0, \sigma_\beta^2)$	σ_β^2	σ_β^2	$\sigma_\beta^2 = \frac{h^2 \sigma_p^2}{\text{MS}_X}$
Bayesian LASSO			
$\text{DE}(\beta_j \sigma^2, \lambda^2)$	$\{\sigma^2, \lambda^2\}$	$2 \frac{\sigma^2}{\lambda^2}$	$\lambda = \sqrt{2 \frac{(1-h^2)}{h^2}} \text{MS}_X$
BayesA			
$t(\beta_j \text{d.f.}_\beta, S_\beta)$	$\{\text{d.f.}_\beta, S_\beta\}$	$\frac{\text{d.f.}_\beta S_\beta^2}{\text{d.f.}_\beta - 2}$	$S_\beta^2 = \frac{(\text{d.f.}_\beta - 2) h^2 \sigma_p^2}{\text{d.f.}_\beta \text{MS}_X}$
Spike-slab			
$\pi \times N\left(\beta_j 0, \frac{\sigma_\beta^2}{\tau}\right) + (1-\pi)N(\beta_j 0, \sigma_\beta^2),$ ($\tau > 1$)	$\{\pi, \sigma_\beta^2, \tau\}$	$\sigma_\beta^2 \times \left[1 + \pi \frac{(1-\tau)}{\tau}\right]$	$\sigma_\beta^2 = \left[\frac{\tau}{\tau + \pi(1-\tau)}\right] \frac{h^2 \sigma_p^2}{\text{MS}_X}$
BayesC			
$\pi \times 1(\beta_j = 0) + (1-\pi)N(\beta_j 0, \sigma_\beta^2)$	$\{\pi, \sigma_\beta^2\}$	$\sigma_\beta^2 \times (1-\pi)$	$\sigma_\beta^2 = \frac{1}{(1-\pi)} \frac{h^2 \sigma_p^2}{\text{MS}_X}$
BayesB			
$\pi \times 1(\beta_j = 0) + (1-\pi)t(\beta_j \text{d.f.}_\beta, S_\beta)$	$\{\pi, \text{d.f.}_\beta, S_\beta\}$	$(1-\pi) \frac{\text{d.f.}_\beta S_\beta^2}{\text{d.f.}_\beta - 2}$	$S_\beta^2 = \frac{1}{(1-\pi)} \frac{(\text{d.f.}_\beta - 2) h^2 \sigma_p^2}{\text{d.f.}_\beta \text{MS}_X}$

$\text{MS}_X = n^{-1} \sum_{i=1}^n \sum_{j=1}^p (x_{ij} - \bar{x}_j)^2$ where $x_{ij} \in (0, 1, 2)$ represents number of copies of the allele coded as one at the j^{th} ($j = 1, \dots, p$) locus of the i^{th} ($i = 1, \dots, n$) individual, and \bar{x}_j is the average genotype at the j^{th} marker.

assigned to marker effects. Examples of such information include, but are not limited to, (a) location of the marker in the genome; (b) whether the marker is located in a coding or a noncoding region; (c) whether the marker is in a region that harbors genes that we may believe affect the trait of interest; and (d) any prior information about the marker from an independent study, such as P -values or estimates of effects derived from a genome-wide association study. One of the great advantages of the Bayesian framework is that we can potentially include all these different types of information into the model via prior specification. With increasing volumes of information coming from multiple studies, the topic of how prior information can be incorporated into models is becoming increasingly important.

One possibility is to assign different priors for different sets of markers, an approach referred to as variance decomposition or variance partition. For instance, in BRR or in G-BLUP one can estimate variance parameters peculiar to chromosomes (Calus *et al.* 2010). Such an approach could also be used with any other prior information that allows grouping the markers such as information about gene function or ontology. Another possibility is to structure the prior to induce borrowing of information across marker effects. For instance, Yang and Tempelman (2012) propose using an antedependence covariance function to specify prior covariances between marker effects.

Algorithms

In Bayesian analysis inferences are based on the posterior distribution of the unknowns given the data with the general form of the posterior density is given in (3). In

most cases, especially when prior hyperparameters are treated as random, the posterior distribution does not have a closed form. However, features of the posterior distribution (e.g., the posterior mean or standard deviation of marker effects) can be approximated using Monte Carlo Markov chain (MCMC) methods (Gelman *et al.* 2003).

Gibbs sampler: Among the many MCMC algorithms the Gibbs sampler (Geman and Geman 1984; Casella and George 1992) is the most commonly used. In a Gibbs sampler draws from the joint posterior density are obtained by sampling from fully conditional densities; therefore, this algorithm is convenient when the fully conditional densities have closed form and are easy to sample from. This occurs, for example, in BRR, where all fully conditionals have closed form. However, this does not directly occur when the priors assigned to marker effects are from the thick-tailed family. To circumvent this problem, the most common approach consists of representing the thick-tailed densities as mixtures of scaled-normal densities (see *Thick-tailed priors* section above). With this approach, used in models BayesA, BayesB, and BL, the fully conditional densities of marker effects as well as those of the conditional variances of marker effects have closed forms. The typical iterations of the Gibbs sampler are illustrated next, using model BayesA as an example.

1. Update the intercept with a sample drawn from a normal density with mean equal to $n^{-1} \sum_{i=1}^n \tilde{y}_i$ and variance $\sigma^2 n^{-1}$, where $\tilde{y}_i = y_i - \sum_{j=1}^p x_{ij} \beta_j$.
2. For j in $\{1, \dots, p\}$ update marker effects with a draw from a normal density with mean $[\sum_{i=1}^n x_{ij}^2 + \sigma^2 \sigma_{\beta_j}^{-2}]^{-1} \sum_{i=1}^n x_{ij} \tilde{y}_i$ and variance $\sigma^2 [\sum_{i=1}^n x_{ij}^2 + \sigma^2 \sigma_{\beta_j}^{-2}]^{-1}$, where

$\tilde{y}_i = y_i - \mu - \sum_{k \neq j} x_{ik} \beta_k$ and $\sigma_{\beta_j}^2$ is the prior conditional variance of the j th marker effect.

3. For j in $\{1, \dots, p\}$ update the variance of marker effects with a draw from a scaled-inverse chi-square density with scale and degree of freedom parameters $\beta_j^2 + \text{d.f.}_\beta \times S_\beta$ and $1 + \text{d.f.}_\beta$, respectively, where S_β and d.f._β are the prior scale and prior degree of freedom assigned to the variances of marker effects.
4. Update the residual variance, σ^2 with a draw from a scaled-inverse chi-square density with degree of freedom $n + \text{d.f.}_e$ and scale $\sum_{i=1}^n \varepsilon_i^2 + \text{d.f.}_e \times S_e$. Here, S_e and d.f._e are the prior scale and prior degree of freedom assigned to the residual variance and $\{\varepsilon_i = y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j\}$ are the current model residuals.

Updating location parameters (intercept and marker effects) requires forming offsets obtained by subtracting from the phenotypes all the regression terms except the one that is being updated. In practice it is computationally less demanding to form the same offset by adding to the current residuals the current sample of the effect that will be updated. For instance, the offset required for sampling the j th marker effect $\tilde{y}_i = y_i - \mu - \sum_{k \neq j} x_{ik} \beta_k$ can be also formed by adding to the current residual the contribution to the regression of the marker whose effect will be updated; that is, $\tilde{y}_i = y_i - \mu - \sum_{k \neq j} x_{ik} \beta_k = \varepsilon_i + x_{ij} \beta_j$. Once location parameters are updated, residuals are updated by subtracting from the offset the contribution to the conditional expectation of the effect just updated (e.g., after drawing the j th marker effect, the updated residuals are $\varepsilon_i = \tilde{y}_i - x_{ij} \beta_j$). Steps 3 and 4 require updating dispersion parameters. In most models the residual variance is updated from an inverse chi-square density. The fully conditional density of the variances of marker effects changes across models. In BayesA these are also inverse chi square. Note that these densities depart from the prior density by only 1 d.f. (see step 3, above), suggesting that data contain little information about these unknowns and that the influence of the prior on inferences about these unknowns can be substantial (Gianola *et al.* 2009).

The structure of the Gibbs sampler for BRR and BL is very similar to that described in steps 1–4, above. For BL the structure is very similar with two main differences: $\text{Var}(\beta_j) = \tau_j^2 \sigma^2$ and τ_j^2 has an exponential prior. Therefore, in step 3 the τ_j^2 are updated from inverse Gaussian densities, and in step 4, $S = \sum_{j=1}^p \tau_j^{-2} \beta_j^2 + \sum_{i=1}^n \varepsilon_i^2 + \text{d.f.}_e \times S_e$ and $\text{d.f.} = p + n + \text{d.f.}_e$. For BRR, $\sigma_{\beta_j}^2 = \sigma_\beta^2$ is the same for all markers; therefore: (a) in step 2 (see above), the term $\sigma^2 \sigma_{\beta_j}^{-2} = \sigma^2 \sigma_\beta^{-2}$ is also the same for all markers and (b) in step 3 only one variance parameter is updated, in this case from a scaled-inverse chi-square density with scale and degree of freedom parameters $\sum_{j=1}^p \beta_j^2 + \text{d.f.}_\beta \times S_\beta$ and $p + \text{d.f.}_\beta$, respectively. Therefore, in this case the fully conditional density departs from the prior by p d.f.

Although the Gibbs sampler is extremely flexible and in general easy to implement, the computational burden increases linearly with the number of records (due to computation of

offsets) and with the number of markers (see steps 2 and 3 above). Many future applications will be using large data sets with hundreds of thousands of markers. In this context the Gibbs sampler can be extremely computationally demanding. To circumvent this problem, a few “fast” methods have been developed; these are briefly discussed next.

Fast methods attempt to estimate the posterior mode by maximizing the posterior density using expectation-maximization (Dempster *et al.* 1977) type algorithms (Yi and Banerjee 2009; Hayashi and Iwata 2010; Shepherd *et al.* 2010). Other proposals attempt to approximate the posterior mean using iterative-conditional expectation procedures (Meuwissen *et al.* 2009). Finally, a third approach consists of using two-step procedures where first, parameters other than marker effects are estimated from their marginal posterior density, and subsequently marker effects are estimated conditional on the other parameters (Cai *et al.* 2011). Thus far, only few studies have compared any of those methods with their MCMC-based counterparts and generally concluded that accuracies of the fast implementations were close to those of the MCMC counterparts; however, it is important to note that many of the algorithms are heuristic and the convergence properties are not well known. This is particularly relevant because in many of the models used in GS the posterior density is not guaranteed to be unimodal; therefore, there is great risk for the algorithm to arrive at, and not move from, local maxima. Moreover, unlike MCMC algorithms, the fast methods usually do not provide estimates of uncertainty about the estimated marker effects or predicted breeding values.

Two-step approaches for G-BLUP: Above we discussed algorithms that estimate variance parameters and marker effects simultaneously. In G-BLUP there are only two variance parameters to be estimated, σ^2 and σ_u^2 . All data points contribute information about these unknowns; therefore with moderate to large sample size and with genetically related individuals the posterior densities of these unknowns are reasonably sharp. Consequently, a common approach consists of first estimating these variance parameters using a non-Bayesian algorithm, usually restricted maximum likelihood, and subsequently computing BLUP of genetic effects from standard mixed-model equations (see Equation 4). This is computationally much more convenient than the MCMC algorithms described above. The form of the restricted maximum-likelihood (REML) objective function and that of the mixed-model equations of G-BLUP are the same as those used in standard pedigree-based models with **G** replacing **A**. However, unlike **G**, which is usually a dense matrix, **A** and its inverse are sparse and most available software developed for pedigree models use sparse-matrix algorithms. Therefore, despite the similarity between G-BLUP and standard BLUP using **A**, some of the existing BLUP software cannot be readily used for G-BLUP. However, several packages such as ASReml (Gilmour *et al.* 2009) and DMU (Madsen and Jensen 2010) have

options that allow providing a dense matrix, \mathbf{G} , or its inverse, instead of pedigree data.

Selected Topics Emerging in Empirical Applications

The application of models for GS to real data usually involves several preprocessing steps. These steps can, in some cases, have great impact on model performance. Here we discuss selected topics that emerge in application of GS with real data.

Coding of marker genotypes

Centering and standardization of covariates is common practice in regression. When regression coefficients are estimated using OLS, centering and standardizing predictors have no effect on predictions. However, when shrinkage estimation procedures are used, transformations of the predictors can potentially impact estimates of effects and predictions. Centering is less relevant because models include an intercept that is usually not penalized; therefore, the effects of centering are “absorbed” in the intercept (Strandén and Christensen 2011). However, rescaling genotypes does have an effect. We explain this from a Bayesian perspective. Let $\tilde{x}_{ij}\tilde{\beta}_j = (x_{ij}/\sqrt{2\theta_j(1-\theta_j)})\tilde{\beta}_j$ be the contribution to the genomic value of the j th marker when genotypes are standardized to a unit variance. Here $x_{ij} \in \{0, 1, 2\}$ are genotype codes in original scale, θ_j is the frequency of the allele coded as one, and $\tilde{\beta}_j$ is the effect of the j th marker when genotypes are standardized. Further, following standard assumptions (Equation 2) let $\sigma_{\tilde{\beta}}^2$ denote the prior variance of the effects of standardized markers. It follows that the effects of unstandardized markers are $\beta_j = \tilde{\beta}_j/\sqrt{2\theta_j(1-\theta_j)}$ ($j = 1, \dots, p$). From here we observe that $\text{Var}(\beta_j) = \sigma_{\tilde{\beta}}^2/\sqrt{2\theta_j(1-\theta_j)}$. The denominator in the expression is maximum at intermediate allele frequencies and approaches zero as θ_j approaches either zero or one; therefore, assigning equal prior variances for the effects of standardized markers implies smaller prior variance (*i.e.*, strongest shrinkage toward zero) for the effects of markers with intermediate allele frequencies and less informative priors (*i.e.*, larger prior variance) for the effects of markers with extreme allele frequencies. In short, other things being equal, standardization induces less (more) shrinkage of estimates of markers with extreme (intermediate) allele frequency.

Preadjusting phenotypes with estimates of systematic effects

The models discussed so far ignore effects other than those of markers and the intercept. In practice, phenotypes are affected by nongenetic effects such as those of contemporary groups (*e.g.*, herd-year-season, sex, or age in animals). Theoretically, one can extend the model described in expression (2) by adding effects other than the intercept and markers. This allows joint estimation of all effects and when this is feasible, joint estimation of effects is the preferred option. However, in practice the joint estimation of effects

may not be feasible because the software available for GS does not allow that or because of high computational requirements or because raw data cannot be shared. In such instances, a common approach is to preadjust data with estimates of nongenetic effects. However, caution must be exercised because, for reasons discussed below, precorrection of data can have undesirable consequences.

Bias: In practice, marker effects and nonmarker covariates are not orthogonal to each other; therefore, estimation in two steps is likely to induce bias (and inconsistency) of estimates of marker effects. For instance, this may occur if genetically superior individuals are overrepresented in some contemporary groups. Similar problems emerge if selection takes place and data are preadjusted with estimates of year effects. All these problems are mitigated if the size of the contemporary groups is large and when the assignment of individuals to groups is at random. In an attempt to reduce biases induced by precorrection, a common practice is to add into the model for precorrection a genetic effect, commonly a pedigree regression. Preadjusted records are then computed by summing predictions of genetic and residual effects. This approach may reduce some of the confounding between genetic and nongenetic effects but it can also bring other types of biases. For instance, in comparing pedigree vs. marker-based models, precorrections based on a pedigree model may “favor” the pedigree model in the second step.

Induced heterogeneous residual variances and residual correlation: In most cases, preadjustments are carried out using a linear model; therefore predictions and fitted residuals are linear functions of data of the form $\hat{\mathbf{y}} = \mathbf{H}\mathbf{y}$ and $\hat{\boldsymbol{\varepsilon}} = \mathbf{y} - \hat{\mathbf{y}} = (\mathbf{I} - \mathbf{H})\mathbf{y}$, respectively, where: $\hat{\mathbf{y}}$ and $\hat{\boldsymbol{\varepsilon}}$ are predicted phenotypes and model residuals, and \mathbf{H} is the so-called hat matrix. For example, if the model for precorrection has the form $\mathbf{y} = \mathbf{W}\mathbf{b} + \boldsymbol{\varepsilon}$ and \mathbf{b} is estimated by OLS, then $\mathbf{H} = \mathbf{W}[\mathbf{W}'\mathbf{W}]^{-1}\mathbf{W}'$. The variance-covariance matrix of the adjusted residuals is $\text{Cov}(\hat{\boldsymbol{\varepsilon}}) = (\mathbf{I} - \mathbf{H})\text{Cov}(\mathbf{y})(\mathbf{I} - \mathbf{H})' = \boldsymbol{\Sigma}$. This matrix is likely to have heterogeneous entries in the diagonal and nonzero off diagonals, implying that predicted residuals may have heterogeneous variances and correlations induced by the preprocessing. A common practice is to “weight” the second-step regression by dividing data with square roots of the diagonal elements of $\boldsymbol{\Sigma}$. Such scaling makes the residual variance homoskedastic; however, this does not account for correlations between residuals that may have been induced by preadjustment of the data.

Dealing with unphenotyped and ungenotyped individuals

In many instances the set of individuals for which genotypic records are available may be different from that of the individuals with genotypes; for example, in dairy cattle many productive and reproductive traits are measured in females and, most commonly, a large proportion of

genotypes are from sires. Schematically, we have two data sets: one comprising phenotypic records and pedigree and one comprising genotypic records. These two sets may be completely disjoint or may partially overlap. Roughly speaking, there are two main strategies that can be used to deal with unphenotyped and ungenotyped individuals.

Two-step genomic evaluations: A common approach is to preprocess the phenotypic data in a way that a phenotypic score (\ddot{y}_i) is produced for each of the genotyped individuals. Examples of these are the use of so-called daughter-yield deviations (DYD), predicted transmitting abilities (PTA), or deregressed proofs (DP) as “phenotypes” in genomic models (VanRaden and Wiggans 1991; Garrick *et al.* 2009). For reasons similar to those discussed in the previous section, these procedures are likely to induce not only heterogeneous residual variances but also correlations between the phenotypic scores, \ddot{y}_i , unless special care is taken in computing these scores (Garrick *et al.* 2009). Heterogeneous residual variances can easily be accounted for in the following genomic evaluation by weighing the residual variance for each individual with an appropriate weight (Fikse and Banos 2001). And Garrick *et al.* (2009) discuss alternatives for computing family means in ways that avoid inducing residual correlations.

Single-step evaluations: The single-step evaluations aim at combining information from genotyped and ungenotyped individuals in a single analysis. This requires deriving the joint distribution of the genetic effects of ungenotyped (\mathbf{g}_1) and genotyped (\mathbf{g}_2) animals, and most proposals attempt to do this in G-BLUP type models. We show in [File S1](#) that if the genotypes of some individuals are unobserved, the joint density of $(\mathbf{g}_1, \mathbf{g}_2)$ is not multivariate normal; rather, it is a mixture of multivariate normal densities (see [File S1](#)). Therefore, existing proposals for combining data from genotyped and ungenotyped individuals (Legarra *et al.* 2009; Christensen and Lund 2010) that assume multivariate normality should be regarded as linear approximations to a non-linear problem.

In the proposed methods (Aguilar *et al.* 2010; Christensen and Lund 2010) the standard genomic relationship matrix, \mathbf{G} , is replaced with a matrix, $\bar{\mathbf{G}}$, computed using observed genotypes (from the subset of genotyped individuals) and pedigree information (which is assumed to include all individuals with phenotypes or genotypes or both sources of information). In $\bar{\mathbf{G}}$, genotypic information propagates into the relationships of ungenotyped individuals, using a linear regression procedure that implicitly predicts unobserved genotypes as linear combinations of observed genotypes with regression coefficients derived from pedigree relationships. The inverse of $\bar{\mathbf{G}}$, which can be used to compute G-BLUP (see Equation 4), has a relatively simple form (Aguilar *et al.* 2010; Christensen and Lund 2010); however, computing $\bar{\mathbf{G}}^{-1}$ requires inverting the matrix of genomic relationships of individuals with genotypes that may be singular. To

circumvent this problem several procedures have been proposed (see Aguilar *et al.* 2010).

Lessons Learned from Simulation and Empirical Data Analysis

In the past few years, many studies have evaluated the performance of various WGP methods. In this section, we review this literature with a focus on extracting what we consider are the lessons learned after almost a decade of research in GS. Early publications were mainly based on simulated data but in recent years empirical evaluations have become more important. Based on a review of published articles we present in Table 2 a list of methods whose predictive performance has been compared to at least one other method for GS. In addition to the method's name and abbreviation we indicate whether the estimation procedure is a penalized or a Bayesian regression. Most of the methods are parametric, in the sense that genomic values are represented as parametric functions of marker genotypes (see Equation 2), but some are nonparametric in nature and this is indicated in the last column of Table 2.

Using the methods listed in Table 2 and a sample of articles we reviewed (references provided at the bottom of Table 2), we provide in Figure 3 the number of times each pair of methods was compared, with the diagonal entries giving the number of times a particular method was used in a comparison study. Figure 3A counts simulation-based studies and Figure 3B summarizes those based on real data. The great majority of the studies included in our review evaluated linear regressions. Nonparametric procedures have the potential of capturing nonadditive effects as well; however, the use of nonparametric procedures remains limited thus far.

Among the additive models, the Bayesian regressions and G-BLUP were the most commonly used. G-BLUP is attractive because its implementation is straightforward using existing REML+BLUP software. The Bayesian methods have been discussed and described widely and are generally chosen because many of them (e.g., BayesA, -B, and -C and BL) allow departures from the infinitesimal model. Many of the penalized regressions (e.g., LASSO and EN) also allow this, but their use in GS is much more limited. This may appear to be surprising considering the diversity of available methods [RR, partial least-squares regression (PLS), principal component regression (PCR), LASSO, EN, and support vector regression (SVR)] and the fact that most of these methods are implemented in relatively efficient packages. However, although limited, empirical evidence suggests somewhat lower prediction accuracy of these methods relative to the more frequently used Bayesian regressions (Solberg *et al.* 2009; Coster *et al.* 2010; Gredler *et al.* 2010; Pszczola *et al.* 2011; Heslot *et al.* 2012).

Simulation studies

Simulation studies (Meuwissen *et al.* 2001; Habier *et al.* 2007) have systematically shown higher prediction accuracy

Table 2 Classification and abbreviations of the models included in Figure 3, A and B

Name (abbreviation)	Bayesian	Penalized	Nonparametric
Least-squares regression (LSR)			
Bayesian ridge regression (BRR) or RR-BLUP	X	X	
BLUP using a genomic relationship matrix (G-BLUP)	X	X	
Trait-specific BLUP (TA-BLUP)	X	X	
BayesA	X		
BayesB	X		
BayesC	X		
Bayes SSVS	X		
Bayesian LASSO (BL)	X		
Double hierarchical generalized linear models (DHGLM)			
Least absolute shrinkage and selection operator (LASSO)		X	
Partial least-squares regression (PLS)		X	
Principal component regression (PCR)		X	
Elastic net (EN)		X	
Reproducing kernel Hilbert spaces regressions (RKHS)	X	X	X
Support vector regression (SVR)		X	X
Boosting ^a	NA	NA	NA
Random forests (RF)			X
Neural networks (NN) ^b	X	X	X

The following are early references of the use of the above methods for genomic prediction (references with the original description of some of the methods are also given in earlier sections of this article and in the references given here). LSR, BRR, BayesA, and BayesB, Meuwissen *et al.* (2001); G-BLUP, VanRaden (2008); TA-BLUP, Zhang *et al.* (2010); BayesC, Habier *et al.* (2011); Bayes SSVS, Calus *et al.* (2008); BL, de los Campos *et al.* (2009); DHGLM, Shen *et al.* (2011); LASSO, Usai *et al.* (2009); PLS and SVR, Moser *et al.* (2009); PCR, Solberg *et al.* (2009); EN, croiseau *et al.* (2011); RKHS, Gianola *et al.* (2006); Boosting, González-Recio *et al.* (2010); RF, González-Recio and Forni (2011); and NN, Okut *et al.* (2011).

^a Boosting as an estimation technique could be applied to any method, Bayesian or penalized, parametric or nonparametric.

^b NN could be implemented in a nonpenalized, penalized, or Bayesian framework.

of GS relative to standard pedigree-based predictions regardless of trait architecture or model of choice. However, many studies have shown that factors such as size of the reference data set (Meuwissen *et al.* 2001; VanRaden and Sullivan 2010), trait heritability, the number of loci affecting the trait (Daetwyler *et al.* 2008), and the degree of genetic relationships between training and validation samples (Habier *et al.* 2007) can greatly affect the prediction accuracy of WGP.

Effects of genetic architecture, marker density, and model on prediction accuracy: The choice of model has also been shown to affect predictive performance, but simulation studies suggest that the effect of model choice on prediction accuracy depends on genetic architecture.

Daetwyler *et al.* (2010b) showed in a simulation study that the accuracy of G-BLUP is not affected by the number of QTL; however, the predictive performance of BayesC was greatly affected by that factor and its accuracy was higher when the number of QTL was low and decreased with increasing number of QTL. Similar trends were observed by Coster *et al.* (2010) and Clark *et al.* (2011). In the study of Coster *et al.* (2010), variable selection methods (*e.g.*, some Bayesian regressions and LASSO) showed an increase in accuracy when the number of QTL decreased, while a method that includes all SNPs (PLS) was shown to be unaffected by the number of QTL. In the study by Clark *et al.* (2011) it was also shown that with BayesB greater accuracies were obtained than with G-BLUP in scenarios involving different distributions of allele frequencies (from rare to common variants) at casual loci. For a scenario that closely resembled the infinitesimal model there was no dif-

ference in accuracy between BayesB and GBLUP, while Daetwyler *et al.* (2010b) reported a lower accuracy for BayesC in such a scenario. Other comparisons based on simulated data also show that there are differences in accuracies across models and generally confirm that, as one would expect, accuracy is greater for the model that better fits the genetic architecture of the trait (Lund *et al.* 2009; Bastiaansen *et al.* 2010; Pszczola *et al.* 2011).

Apart from the number of QTL and the distribution of their effects, there are several other factors that theoretically are expected to result in differences in accuracy for variable selection methods vs. G-BLUP type models. With low marker density markers are unlikely to be tightly linked to QTL and each marker may track signal from different QTL, inducing a less extreme distribution of effects than obtained at higher density. Meuwissen *et al.* (2009) showed that the superiority of BayesB over G-BLUP increased with increasing marker density. This result is also clearly confirmed in the study by Meuwissen and Goddard (2010), where genomic prediction based on the whole-genome sequence including the causal loci was substantially more accurate for BayesB compared to G-BLUP. However, both studies simulated very few QTL, giving variable selection methods an advantage. On the other hand, for any given marker density, the ability of variable selection methods to detect variants tightly linked to QTL increases as the span of LD decreases.

Real data analysis

A number of empirical evaluations of GS have been published in recent years. These studies have confirmed some but not all of the findings anticipated by simulation studies.




Figure 3 (A and B) Number of articles reviewed comparing one or more methods using simulated (A) or real (B) data. The abbreviations used for the methods are given in Table 2. The following references were used: (Meuwissen *et al.* 2001; Habier *et al.* 2007; Piyasatian *et al.* 2007; González-Recio *et al.* 2008; Lee *et al.* 2008; Bennewitz *et al.* 2009; de los Campos *et al.* 2009; Gonzalez-Recio *et al.* 2009; Hayes *et al.* 2009a,b; Lorenzana and Bernardo 2009; Luan *et al.* 2009; Lund *et al.* 2009; Meuwissen 2009; Meuwissen *et al.* 2009; Moser *et al.* 2009; Solberg *et al.* 2009; Usai *et al.* 2009; Verbyla *et al.* 2009; Zhong *et al.* 2009; Andreescu *et al.* 2010; Bastiaansen *et al.* 2010; Coster *et al.* 2010; Crossa *et al.* 2010; Daetwyler *et al.* 2010a,b; de los Campos *et al.* 2010a,b; Gonzalez-Recio *et al.* 2010; Gredler *et al.* 2010; Guo *et al.* 2010; Habier *et al.* 2010; Konstantinov and Hayes 2010; Meuwissen and Goddard 2010; Mrode *et al.* 2010; Pérez *et al.* 2010; Shepherd *et al.* 2010; Zhang *et al.* 2010; Calus and Veerkamp 2011; Clark *et al.* 2011; Croiseau *et al.* 2011; de Roos *et al.* 2011; Gonzalez-Recio and Forni 2011; Habier *et al.* 2011; Heffner *et al.* 2011; Iwata and Jannink 2011; Legarra *et al.* 2011; Long *et al.* 2011a,b; Makowsky *et al.* 2011; Mujibi *et al.* 2011; Ober *et al.* 2011; Ostensen *et al.* 2011; Pryce *et al.* 2011; Pszczola *et al.* 2011; Wiggans *et al.* 2011; Wittenburg *et al.* 2011; Wolc *et al.* 2011a,b; Yu and Meuwissen 2011; Bastiaansen *et al.* 2012; Heslot *et al.* 2012).

Pedigree vs. marker-enabled prediction: In general, empirical studies have confirmed the superiority of GS relative to family-based predictions that was anticipated by simulations. The most clear case occurs in Holstein dairy cattle where several studies (Hayes *et al.* 2009a; VanRaden *et al.* 2009) have confirmed that GS can attain higher prediction accuracy than standard family-based prediction (e.g., parental average). The potential of GS has also been confirmed in several breeds of beef cattle (Garrick 2011; Saatchi *et al.* 2011), sheep (Daetwyler *et al.* 2010a), broilers (Gonzalez-Recio *et al.* 2008), layer chickens (Wolc *et al.* 2011a,b), and several plant species (de los Campos *et al.* 2009; Crossa *et al.* 2010; Heslot *et al.* 2012). For applications in plants, it has been shown that GS outperforms conventional marker-assisted selection (Heffner *et al.* 2010, 2011) and that it has the potential to be substantially more efficient per unit of time than phenotypic selection (Grattapaglia and Resende 2011; Zhao *et al.* 2012). However, the superiority of GS, relative to pedigree-based predictions, has not always been as high as anticipated by simulations. This is particularly clear in applications for sheep (Daetwyler *et al.* 2010a) and in beef cattle (Saatchi *et al.* 2011). Many reasons may contribute to this: in general in these cases the size of the training data set is

limited and the phenotypes used may have been noisy, in some cases (e.g., some breeds in sheep) population structure (e.g., a mixture of breeds) may be a reason, and in others the relatively small family size may limit the potential accuracy that can be attained with GS.

Some studies have shown benefits of extending the standard models of GS by adding a random effect representing a regression on pedigree information (de los Campos *et al.* 2009; Crossa *et al.* 2010); however, the benefits of jointly modeling pedigree and marker data relative to a markers-only model seem to vanish as marker density increases (Vazquez *et al.* 2010).

Effects of marker density: Several empirical studies have evaluated the effects of marker density on prediction accuracy (Weigel *et al.* 2009; Vazquez *et al.* 2010; Makowsky *et al.* 2011). When shrinkage estimation procedures are used, prediction accuracy increases monotonically with marker density, but it does at diminishing marginal rates of response. Consequently, prediction accuracy reaches a plateau and does not increase beyond certain marker density. The level at which this plateau takes place depends on two factors mainly: the span of LD in the genome and sample

size. For instance, Vazquez *et al.* (2010) found little increase on prediction accuracy beyond 10,000 SNPs for the prediction of several traits in U.S. Holsteins, a population where LD span over long regions. However, in a study involving WGP of human data, which exhibit much shorter span of LD than found in cattle data, Makowsky *et al.* (2011) found response to increased marker density even beyond 100,000 SNPs; more importantly the rates of response to increases in marker density were greatly affected by the number of close relatives used in the training data set, suggesting that “local sample size” also affects the effects of marker density on prediction accuracy.

Genetic architecture, sample size, and model: For traits involving a limited number of large-effect QTL, simulation studies have consistently predicted superiority of methods using variable selection and differential shrinkage of estimates of effects such as BayesB. However, this has not been fully confirmed by real data analysis. Indeed, in most studies comparing different genomic prediction models based on real data, there are only small differences observed in accuracies between models. An important question is whether the genetic architecture of real data is perhaps less extreme than suggested by QTL-mapping studies (Kearsey and Farquhar 1998; Hayes and Goddard 2001) and generally assumed in simulation studies or whether other characteristics of the data (number of markers, span of LD) used prohibit greater distinction between models. Variable selection is most effective with short span of LD, high marker density, and large sample size; however, these conditions are not always met in applications in plant and animal breeding.

In empirical studies genetic architecture is less well known and grouping traits based on their architecture is not straightforward. In animal breeding, there are only a few examples of traits where one or a few major genes explain a sizable proportion of genetic variance. One of such examples is the *DGAT1* gene that has a large effect on fat percentage in dairy cattle (Grisart *et al.* 2002; Winter *et al.* 2002). For this particular case, it is shown in several studies that models with a thick-tailed prior distribution of marker effects such as BayesA and variable selection methods such as Bayes SSVS yield higher accuracy than G-BLUP.

In Figure 4 we summarize results from three studies (Hayes *et al.* 2009b; Verbyla *et al.* 2009; de Roos *et al.* 2011) where prediction accuracy of G-BLUP and of a Bayesian regression using either a thick-tailed (BayesA) or a spike-slab (Bayes SSVS) density was evaluated for fat and protein percentage at varying sizes of the training data set. The results of these studies illustrate important concepts. First, prediction accuracy increases markedly as the size of the training data set does. This has been anticipated by simulation studies (VanRaden and Sullivan 2010) and consistently confirmed in empirical analyses (Lorenzana and Bernardo 2009; Bastiaansen *et al.* 2010). Second, in general, for fat percentage there is a clear superiority of the models performing differential shrinkage of estimates of effects (*e.g.*, Bayes SSVS or BayesA) relative to G-

BLUP. All traits were analyzed with model BayesA and the authors concluded that prediction accuracy was higher for traits with simpler genetic architecture. This also represents a confirmation of results from simulations that anticipated superiority of methods performing variable selection and differential shrinkage of estimates, especially with traits having some large-effects QTL (Meuwissen *et al.* 2001; Daetwyler *et al.* 2010b). However, the differences detected in empirical evaluations are not as large as those usually reported in simulations. More importantly, differences in predictive performance between methods decreased when the sample size increased (see Figure 4), reflecting that, as expected, the influence of the prior (the choice of prior density of marker effects in this case) decreases as sample size increases. This is in agreement with findings from simulation studies such as that of Daetwyler *et al.* (2010b).

A recent study on genomic prediction for different traits in loblolly pine indicated generally small differences in prediction accuracy between the models RR-BLUP, BL, BayesA, and BayesC (Resende *et al.* 2012). Nevertheless, for traits related to fusiform rust, known to be controlled by a few genes of large effects, BayesA and BayesC outperformed RR-BLUP and BL.

Theoretically, it can be expected that a method that is flexible enough to model any type of genetic architecture (*e.g.*, model BayesB or BayesC when all hyperparameters are estimated from data) will perform relatively well across a wide range of traits with different characteristics; however, empirical evidence has not always confirmed this. For instance, the study of Heslot *et al.* (2012) compared 10 different methods for 18 different traits measured on barley, wheat, and maize, and RKHS was the only method that clearly outperformed others across traits and data sets. Among the linear regressions, the Bayesian methods outperformed the penalized regressions, and the differences among the Bayesian linear regressions were very small. However, marker density in this study was very low and this clearly limits the possibility of some methods to express their potential.

The effects of sample size on prediction accuracy are clear. First, the accuracy of estimates of marker effects increases with sample size. This occurs because bias and variance of estimates of marker effects decrease with sample size. Additionally, in some designs an increase in sample size may also increase the extent of genetic relationships between subjects in the training and validation data sets.

Liu *et al.* (2011) evaluated the effect of the size of the training data set on the dispersion of estimates of marker effects and on prediction accuracy using BRR-BLUP. The authors reported an increase in dispersion among estimates of SNP effects by more than a factor of 5, when the reference population was increased from 735 to 5025 bulls (Liu *et al.* 2011). This essentially reflects that the extent of shrinkage of estimates of effects decreases as sample size increases. Moreover, they showed that the correlation between estimated SNP effects was as low as 0.42 between the two extreme scenarios, but correlations were >0.9 when ~500




Figure 4 Accuracies of G-BLUP, BayesA, and Bayes SSVS models for fat and protein percentage, estimated using three different Holstein–Friesian reference populations (Hayes *et al.* 2009b; Verbyla *et al.* 2009; de Roos *et al.* 2011). Note that the data used by Hayes *et al.* (2009b) are a subset of the data used by Verbyla *et al.* (2009).

bulls were added to the reference population that already contained >3500 bulls. However, they reported that the correlation between genomic predictions with different numbers in the reference population was much closer to 1.0, compared to the correlation between estimated SNP effects. This illustrates two important points: (a) when $p \gg n$, one can arrive at similar predictions of total genomic merit with very different estimates of marker effects and (b), because of the same reason, with small sample size, one needs to be cautious about interpretation of estimates of marker effects because these may be highly influenced by the choice of prior.

Concluding remarks

Both simulation and empirical studies have systematically shown higher prediction accuracy of GS relative to standard pedigree-based predictions and it is now clear that GS offers great opportunities to further increase the rate of genetic progress achieved in plant and animal breeding. Most of the benefits of GS arise from the possibility of obtaining accurate predictions early in the breeding cycle; therefore, getting the most out of GS may require changes in breeding programs.

Implementing GS involves making important decisions regarding the choice of model, the size of the training data, and marker density, just to mention a few. In recent years simulation and empirical studies have produced valuable information that can be used to guide researchers in making those decisions. Several factors, including span of LD, trait heritability, genetic architecture, marker density, size of the training data set, and the model used can affect the prediction accuracy of GS.

Some of the factors affecting prediction accuracy, such as trait heritability, genetic architecture, and to a large extent LD, cannot be controlled; however, we can have control of the design of reference data sets, including size and relationships, marker density, and the model used for estimation of effects. Among the factors that are under control of the researcher, the size of the training data set and the strength of genetic relationships between training and validation samples are by far the most important factors affecting prediction accuracy. The model of choice is also important; however, the differences between models reported by simulation studies have not always been confirmed by real data analysis. Empirical analyses have shown only small differences between methods, with a slight advantage of models performing “selection and shrinkage” such as BayesB for traits with “large-effect QTL.” But in general thick-tailed models such as BayesA or Bayesian LASSO perform well across traits and G-BLUP performs well for most traits. An important reason is that, due to the fact that $p \gg n$, there are a multitude of different prediction equations that yield about the same likelihood and minimum prediction error rate (Breiman 2001): often we encounter multiple equally-good models.

The unexpected generally good performance of G-BLUP in real data is connected to several issues that have been revealed to some extent. First, the real genetic architecture of traits appears less extreme than expected based on QTL-mapping results. Second, most of the gain in accuracy due to using markers in current applications arises from explaining the Mendelian sampling term, rather than from tracing signals generated at individual QTL. Overall, it seems that with a long span of LD and relatively sparse platforms (*e.g.*, 50,000 SNPs) variable selection may not be needed. However, the relative performance of G-BLUP and variable selection methods may change with denser coverage (*e.g.*, with genotyping by sequencing) and in populations with short LD span.

Acknowledgments

The authors thank D. J. de Koning and Lauren McIntyre for encouraging us to write this review article and for insightful comments provided on earlier versions of the manuscript. Mario Calus acknowledges financial support from the Dutch Ministry of Economic Affairs, Agriculture, and Innovation (Programme “Kennisbasis Research,” code KB-17-003.02-006). H.D.D. was funded by the Cooperative Research Centre for Sheep Industry Innovation. J.M.H. was funded by the Australian Research Council project LP100100880 of which Genus Plc, Aviagen LTD, and Pfizer are cofunders.

Note added in proof: See Daetwyler *et al.* 2013 (pp. 347–365) for a related work.

Literature Cited

Aguilar, I., I. Misztal, D. L. Johnson, A. Legarra, S. Tsuruta *et al.*, 2010 Hot topic: a unified approach to utilize phenotypic, full pedigree, and genomic information for genetic evaluation of Holstein final score. *J. Dairy Sci.* 93: 743–752.

- Andresescu, C., D. Habier, R. L. Fernando, A. Kranis, K. Watson *et al.*, 2010 Accuracy of genomic predictions across breeding lines of chickens. 9th World Congress on Genetics Applied to Livestock Production, Leipzig, Germany. CD-ROM Communication 0956.
- Andrews, D. F., and C. L. Mallows, 1974 Scale mixtures of normal distributions. *J. R. Stat. Soc., B* 36: 99–102.
- Bastiaansen, J. W. M., M. C. A. M. Bink, A. Coster, C. Maliepaard, and M. P. L. Calus, 2010 Comparison of analyses of the QTLMAS XIII common dataset. I: Genomic selection. *BMC Proc.* 4(Suppl. 1): S1.
- Bastiaansen, J. W. M., A. Coster, M. P. L. Calus, J. A. M. van Arendonk, and H. Bovenhuis, 2012 Long-term response to genomic selection: effects of estimation method and reference population structure for different genetic architectures. *Genet. Sel. Evol.* 44: 3.
- Bennewitz, J., T. Solberg, and T. Meuwissen, 2009 Genomic breeding value estimation using nonparametric additive regression models. *Genet. Sel. Evol.* 41: 20.
- Bernardo, R., 2008 Molecular markers and selection for complex traits in plants: learning from the last 20 years. *Crop Sci.* 48: 1649–1664.
- Bernardo, R., and J. Yu, 2007 Prospects for genomewide selection for quantitative traits in maize. *Crop Sci.* 47: 1082–1090.
- Breiman, L., 2001 Statistical modeling: the two cultures. *Stat. Sci.* 16: 199–215.
- Buckler, E. S., J. B. Holland, P. J. Bradbury, C. B. Acharya, P. J. Brown *et al.*, 2009 The genetic architecture of maize flowering time. *Science* 325: 714–718.
- Cai, X., A. Huang, and S. Xu, 2011 Fast empirical Bayesian LASSO for multiple quantitative trait locus mapping. *BMC Bioinformatics* 12: 211.
- Calus, M. P. L., and R. F. Veerkamp, 2011 Accuracy of multi-trait genomic selection using different methods. *Genet. Sel. Evol.* 43: 26.
- Calus, M. P. L., T. H. E. Meuwissen, A. P. W. de Roos, and R. F. Veerkamp, 2008 Accuracy of genomic selection using different methods to define haplotypes. *Genetics* 178: 553–561.
- Calus, M. P. L., H. A. Mulder, and R. F. Veerkamp, 2010 Estimation of breeding values for haploid chromosomes. 9th World Congress on Genetics Applied to Livestock Production, Leipzig, Germany. CD-ROM Communication 0269.
- Casella, G., 2001 Empirical Bayes Gibbs sampling. *Biostatistics* 2: 485–500.
- Casella, G., and E. I. George, 1992 Explaining the Gibbs sampler. *Am. Stat.* 46: 167–174.
- Christensen, O. F., and M. S. Lund, 2010 Genomic prediction when some animals are not genotyped. *Genet. Sel. Evol.* 42: 2.
- Clark, S. A., J. M. Hickey, and J. H. J. van der Werf, 2011 Different models of genetic variation and their effect on genomic evaluation. *Genet. Sel. Evol.* 43: 18.
- Collard, B. C. Y., and D. J. Mackill, 2008 Marker-assisted selection: an approach for precision plant breeding in the twenty-first century. *Philos. Trans. R. Soc. B* 363: 557–572.
- Coster, A., J. W. Bastiaansen, M. P. Calus, J. A. van Arendonk, and H. Bovenhuis, 2010 Sensitivity of methods for estimating breeding values using genetic markers to the number of QTL and distribution of QTL variance. *Genet. Sel. Evol.* 42: 9.
- Croiseau, P., A. Legarra, F. Guillaume, S. Fritz, A. Baur *et al.*, 2011 Fine tuning genomic evaluations in dairy cattle through SNP pre-selection with the Elastic-Net algorithm. *Genet. Res.* 93: 409–417.
- Crossa, J., G. de los Campos, P. Perez, D. Gianola, J. Burgueño *et al.*, 2010 Prediction of genetic values of quantitative traits in plant breeding using pedigree and molecular markers. *Genetics* 186: 713–724.
- Daetwyler, H. D., B. Villanueva, and J. A. Woolliams, 2008 Accuracy of predicting the genetic risk of disease using a genome-wide approach. *PLoS ONE* 3: e3395.
- Daetwyler, H. D., J. M. Hickey, J. M. Henshall, S. Dominik, B. Gredler *et al.*, 2010a Accuracy of estimated genomic breeding values for wool and meat traits in a multi-breed sheep population. *Anim. Prod. Sci.* 50: 1004–1010.
- Daetwyler, H. D., R. Pong-Wong, B. Villanueva, and J. A. Woolliams, 2010b The impact of genetic architecture on genome-wide evaluation methods. *Genetics* 185: 1021–1031.
- Daetwyler, H. D., M. P. L. Calus, R. Pong-Wong, G. de los Campos, and J. M. Hickey, 2013 Genomic prediction in animals and plants: simulation of data, validation, reporting, and benchmarking. *Genetics* 193: 347–365.
- Dekkers, J. C., 2004 Commercial application of marker- and gene-assisted selection in livestock: strategies and lessons. *J. Anim. Sci.* 82: E313–E328.
- de los Campos, G., H. Naya, D. Gianola, J. Crossa, A. Legarra *et al.*, 2009 Predicting quantitative traits with regression models for dense molecular markers and pedigree. *Genetics* 182: 375–385.
- de los Campos, G., D. Gianola, G. J. M. Rosa, K. A. Weigel, and J. Crossa, 2010a Semi-parametric genomic-enabled prediction of genetic values using reproducing kernel Hilbert spaces methods. *Genet. Res.* 92: 295–308.
- de los Campos, G., D. Gianola, G. J. M. Rosa, K. Weigel, A. I. Vazquez *et al.*, 2010b Semi-parametric marker-enabled prediction of genetic values using reproducing kernel Hilbert spaces regressions. 9th World Congress on Genetics Applied to Livestock Production, Leipzig, Germany. CD-ROM Communication 0520.
- Dempster, A. P., N. M. Laird, and D. B. Rubin, 1977 Maximum likelihood from incomplete data via the EM algorithm. *J. R. Stat. Soc. B* 39: 1–38.
- de Roos, A. P. W., C. Schrooten, and T. Druet, 2011 Genomic breeding value estimation using genetic markers, inferred ancestral haplotypes, and the genomic relationship matrix. *J. Dairy Sci.* 94: 4708–4714.
- Eding, H., and T. H. E. Meuwissen, 2001 Marker based estimates of between and within population kinships for the conservation of genetic diversity. *J. Anim. Breed. Genet.* 118: 141–159.
- Fernando, R., and M. Grossman, 1989 Marker assisted selection using best linear unbiased prediction. *Genet. Sel. Evol.* 21: 467.
- Fikse, W., and G. Banos, 2001 Weighting factors of sire daughter information in international genetic evaluations. *J. Dairy Sci.* 84: 1759–1767.
- Frank, I. E., and J. H. Friedman, 1993 A statistical view of some chemometrics regression tools. *Technometrics* 35: 109–135.
- Garrick, D., 2011 The nature, scope and impact of genomic prediction in beef cattle in the United States. *Genet. Sel. Evol.* 43: 17.
- Garrick, D. J., J. F. Taylor, and R. L. Fernando, 2009 Deregressing estimated breeding values and weighting information for genomic regression analyses. *Genet. Sel. Evol.* 41: 55.
- Gelman, A., J. B. Carlin, H. S. Stern, and D. B. Rubin, 2003 *Bayesian Data Analysis*, Ed. 2. Chapman & Hall/CRC, London/New York/Washington, D.C./Boca Raton, FL.
- Geman, S., and D. Geman, 1984 Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images. *IEEE Trans. Patt. Anal. Machine Intell.* 6: 721–741.
- George, E. I., and R. E. McCulloch, 1993 Variable selection via Gibbs sampling. *J. Am. Stat. Assoc.* 88: 881–889.
- Gianola, D., R. L. Fernando, and A. Stella, 2006 Genomic-assisted prediction of genetic value with semiparametric procedures. *Genetics* 173: 1761–1776.
- Gianola, D., G. de los Campos, W. G. Hill, E. Manfredi, and R. Fernando, 2009 Additive genetic variability and the Bayesian alphabet. *Genetics* 183: 347–363.
- Gianola, D., G. de los Campos, O. González-Recio, N. Long, H. Okut *et al.*, 2010 Statistical learning methods for genome-based analysis of quantitative traits. *World Genetic Congress Applied*

- to Livestock Production, Leipzig, Germany, CD-ROM Communication 0014.
- Gilmour, A., B. Gogel, B. Cullis, and R. Thompson, 2009 *ASReml User Guide*. VSN International, Hemel Hempstead, UK.
- González-Recio, O., and S. Forni, 2011 Genome-wide prediction of discrete traits using Bayesian regressions and machine learning. *Genet. Sel. Evol.* 43: 1–12.
- Gonzalez-Recio, O., D. Gianola, N. Long, K. A. Weigel, G. J. Rosa *et al.*, 2008 Nonparametric methods for incorporating genomic information into genetic evaluations: an application to mortality in broilers. *Genetics* 178: 2305–2313.
- Gonzalez-Recio, O., D. Gianola, G. J. M. Rosa, K. A. Weigel, and A. Kranis, 2009 Genome-assisted prediction of a quantitative trait measured in parents and progeny: application to food conversion rate in chickens. *Genet. Sel. Evol.* 41: 3.
- González-Recio, O., K. A. Weigel, D. Gianola, H. Naya, and G. J. M. Rosa, 2010 L2-Boosting algorithm applied to high-dimensional problems in genomic selection. *Genet. Res.* 92: 227–237.
- Grattapaglia, D., and M. D. V. Resende, 2011 Genomic selection in forest tree breeding. *Tree Genet. Genomes* 7: 241–255.
- Gredler, B., H. Schwarzenbacher, C. Egger-Danner, C. Fuerst, and R. Emmerling, 2010 Accuracy of genomic selection in dual purpose Fleckvieh cattle using three types of methods and phenotypes. 9th World Congress on Genetics Applied to Livestock Production, Leipzig, Germany. CD-ROM Communication 0907.
- Grisart, B., W. Coppieters, F. Farnir, L. Karim, C. Ford *et al.*, 2002 Positional candidate cloning of a QTL in dairy cattle: identification of a missense mutation in the bovine DGAT1 gene with major effect on milk yield and composition. *Genome Res.* 12: 222–231.
- Guo, G., M. S. Lund, Y. Zhang, and G. Su, 2010 Comparison between genomic predictions using daughter yield deviation and conventional estimated breeding value as response variables. *J. Anim. Breed. Genet.* 127: 423–432.
- Habier, D., R. L. Fernando, and J. M. Dekkers, 2007 The impact of genetic relationship information on genome-assisted breeding values. *Genetics* 177: 2389–2397.
- Habier, D., J. Tetens, F. R. Seefried, P. Lichtner, and G. Thaller, 2010 The impact of genetic relationship information on genomic breeding values in German Holstein cattle. *Genet. Sel. Evol.* 42: 5.
- Habier, D., R. Fernando, K. Kizilkaya, and D. Garrick, 2011 Extension of the Bayesian alphabet for genomic selection. *BMC Bioinformatics* 12: 186.
- Haley, C. S., and P. M. Visscher, 1998 Strategies to utilize marker-quantitative trait loci associations. *J. Dairy Sci.* 81(Suppl. 2): 85–97.
- Hastie, T., R. Tibshirani, and J. Friedman, 2009 *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, Ed. 2. Springer-Verlag, Berlin/Heidelberg, Germany/New York.
- Hayashi, T., and H. Iwata, 2010 EM algorithm for Bayesian estimation of genomic breeding values. *BMC Genet.* 11: 3.
- Hayes, B., and M. E. Goddard, 2001 The distribution of the effects of genes affecting quantitative traits in livestock. *Genet. Sel. Evol.* 33: 209–229.
- Hayes, B. J., P. J. Bowman, A. J. Chamberlain, and M. E. Goddard, 2009a Invited review: genomic selection in dairy cattle: progress and challenges. *J. Dairy Sci.* 92: 433–443.
- Hayes, B., P. Bowman, A. Chamberlain, K. Verbyla, M. Goddard *et al.*, 2009b Accuracy of genomic breeding values in multi-breed dairy cattle populations. *Genet. Sel. Evol.* 41: 51.
- Heffner, E. L., A. J. Lorenz, J.-L. Jannink, and M. E. Sorrells, 2010 Plant breeding with genomic selection: gain per unit time and cost. *Crop Sci.* 50: 1681–1690.
- Heffner, E. L., J.-L. Jannink, H. Iwata, E. Souza, and M. E. Sorrells, 2011 Genomic selection accuracy for grain quality traits in biparental wheat populations. *Crop Sci.* 51: 2597–2606.
- Heslot, N., M. E. Sorrells, J. L. Jannink, and H. P. Yang, 2012 Genomic selection in plant breeding: a comparison of models. *Crop Sci.* 52: 146–160.
- Hoerl, A. E., and R. W. Kennard, 1970 Ridge regression: biased estimation for nonorthogonal problems. *Technometrics* 12: 55–67.
- Hospital, F., 2009 Challenges for effective marker-assisted selection in plants. *Genetica* 136: 303–310.
- Iwata, H., and J. L. Jannink, 2011 Accuracy of genomic selection prediction in barley breeding programs: a simulation study based on the real single nucleotide polymorphism data of barley breeding lines. *Crop Sci.* 51: 1915–1927.
- Kearsey, M. J., and A. G. L. Farquhar, 1998 QTL analysis in plants; where are we now? *Heredity* 80: 137–142.
- Kimeldorf, G. S., and G. Wahba, 1970 A correspondence between Bayesian estimation on stochastic processes and smoothing by splines. *Ann. Math. Stat.* 41: 495–502.
- Konstantinov, K. V., and B. J. Hayes, 2010 Comparison of BLUP and Reproducing kernel Hilbert spaces methods for genomic prediction of breeding values in Australian Holstein Friesian cattle. Proceedings of the 9th World Congress on Genetics Applied to Livestock Production, Leipzig, Germany. CD-ROM Communication 0224.
- Lee, S. H., J. H. J. van der Werf, B. J. Hayes, M. E. Goddard, and P. M. Visscher, 2008 Predicting unobserved phenotypes for complex traits from whole-genome SNP data. *PLoS Genet.* 4: e1000231.
- Legarra, A., I. Aguilar, and I. Misztal, 2009 A relationship matrix including full pedigree and genomic information. *J. Dairy Sci.* 92: 4656–4663.
- Legarra, A., C. Robert-Granie, P. Croiseau, F. Guillaume, and S. Fritz, 2011 Improved Lasso for genomic selection. *Genet. Res.* 93: 77–87.
- Liu, Z., F. R. Seefried, F. Reinhardt, S. Rensing, G. Thaller *et al.*, 2011 Impacts of both reference population size and inclusion of a residual polygenic effect on the accuracy of genomic prediction. *Genet. Sel. Evol.* 43: 19.
- Long, N., D. Gianola, G. J. M. Rosa, and K. A. Weigel, 2011a Dimension reduction and variable selection for genomic selection: application to predicting milk yield in Holsteins. *J. Anim. Breed. Genet.* 128: 247–257.
- Long, N., D. Gianola, G. J. M. Rosa, and K. A. Weigel, 2011b Application of support vector regression to genome-assisted prediction of quantitative traits. *Theor. Appl. Genet.* 123: 1065–1074.
- Lorenzana, R. E., and R. Bernardo, 2009 Accuracy of genotypic value predictions for marker-based selection in biparental plant populations. *Theor. Appl. Genet.* 120: 151–161.
- Luan, T., J. A. Woolliams, S. Lien, M. Kent, M. Svendsen *et al.*, 2009 The accuracy of genomic selection in Norwegian Red cattle assessed by cross-validation. *Genetics* 183: 1119–1126.
- Lund, M. S., G. Sahana, D. J. De Koning, G. Su, and Ö. Carlborg, 2009 Comparison of analyses of the QTLMAS XII common dataset. I: Genomic selection. *BMC Proc.* 3(Suppl. 1): S1.
- Lynch, M., and K. Ritland, 1999 Estimation of pairwise relatedness with molecular markers. *Genetics* 152: 1753–1766.
- Madsen, P. A., and J. Jensen, 2010 *A User's Guide to DMU. A Package for Analysing Multivariate Mixed Models*. Version 6, release 5.0. University of Aarhus, Denmark. http://www.dmu.agrsci.dk/dmuv6_guide-R4-6-7.pdf
- Makowsky, R., N. M. Pajewski, Y. C. Klimentidis, A. I. Vazquez, C. W. Duarte *et al.*, 2011 Beyond missing heritability: prediction of complex traits. *PLoS Genet.* 7: e1002051.
- Meuwissen, T. H. E., 2009 Accuracy of breeding values of 'unrelated' individuals predicted by dense SNP genotyping. *Genet. Sel. Evol.* 41: 35.

- Meuwissen, T., and M. Goddard, 2010 Accurate prediction of genetic values for complex traits by whole-genome resequencing. *Genetics* 185: 623–631.
- Meuwissen, T. H., B. J. Hayes, and M. E. Goddard, 2001 Prediction of total genetic value using genome-wide dense marker maps. *Genetics* 157: 1819–1829.
- Meuwissen, T., T. R. Solberg, R. Shepherd, and J. A. Woolliams, 2009 A fast algorithm for BayesB type of prediction of genome-wide estimates of genetic value. *Genet. Sel. Evol.* 41: 2.
- Moser, G., B. Tier, R. E. Crump, M. S. Khatkar, H. W. Raadsma *et al.*, 2009 A comparison of five methods to predict genomic breeding values of dairy bulls from genome-wide SNP markers. *Genet. Sel. Evol.* 41: 56.
- Mrode, R., M. P. Coffey, I. Strandén, T. H. E. Meuwissen, J. B. C. H. M. Van Kaam *et al.*, 2010 A comparison of various methods for the computation of genomic breeding values of dairy bulls using software at Genomicsselection.net. Proceedings of the 9th World Congress on Genetics Applied to Livestock Production, Leipzig, Germany. CD-ROM Communication 0518.
- Mujibi, F. D. N., J. D. Nkrumah, O. N. Durunna, J. R. Grant, J. Mah *et al.*, 2011 Associations of marker panel scores with feed intake and efficiency traits in beef cattle using preselected single nucleotide polymorphisms. *J. Anim. Sci.* 89: 3362–3371.
- Nadaf, J., V. Riggio, T.-P. Yu, and R. Pong-Wong, 2012 Effect of the prior distribution of SNP effects on the estimation of total breeding value. *BMC Proc.* 6(Suppl. 2): S6.
- Nejati-Javaremi, A., C. Smith, and J. P. Gibson, 1997 Effect of total allelic relationship on accuracy of evaluation and response to selection. *J. Anim. Sci.* 75: 1738–1745.
- Ober, U., M. Erbe, N. Y. Long, E. Porcu, M. Schlather *et al.*, 2011 Predicting genetic values: a kernel-based best linear unbiased prediction with genomic data. *Genetics* 188: 695–708.
- Okut, H., D. Gianola, G. J. M. Rosa, and K. A. Weigel, 2011 Prediction of body mass index in mice using dense molecular markers and a regularized neural network. *Genet. Res.* 93: 189–201.
- Ostersen, T., O. Christensen, M. Henryon, B. Nielsen, G. Su *et al.*, 2011 Deregressed EBV as the response variable yield more reliable genomic predictions than traditional EBV in pure-bred pigs. *Genet. Sel. Evol.* 43: 38.
- Park, T., and G. Casella, 2008 The Bayesian Lasso. *J. Am. Stat. Assoc.* 103: 681–686.
- Pérez, P., G. de los Campos, J. Crossa, and D. Gianola, 2010 Genomic-enabled prediction based on molecular markers and pedigree using the Bayesian Linear Regression package in R. *Plant Gen.* 3: 106–116.
- Piyasatian, N., R. L. Fernando, and J. C. M. Dekkers, 2007 Genomic selection for marker-assisted improvement in line crosses. *Theor. Appl. Genet.* 115: 665–674.
- Pong-Wong, R., A. W. George, J. A. Woolliams, and C. S. Haley, 2001 A simple and rapid method for calculating identity-by-descent matrices using multiple markers. *Genet. Sel. Evol.* 33: 453–471.
- Pryce, J. E., B. Gredler, S. Bolormaa, P. J. Bowman, C. Egger-Danner *et al.*, 2011 Short communication: genomic selection using a multi-breed, across-country reference population. *J. Dairy Sci.* 94: 2625–2630.
- Pszczola, M., T. Strabel, A. Wolc, S. Mucha, and M. Szydlowski, 2011 Comparison of analyses of the QTLMAS XIV common dataset. I: Genomic selection. *BMC Proc.* 5(Suppl. 3): S1.
- Resende, M. F. R. Jr., P. Muñoz, M. D. V. Resende, D. J. Garrick, R. L. Fernando *et al.*, 2012 Accuracy of genomic selection methods in a standard dataset of Loblolly pine (*Pinus taeda* L.). *Genetics* 190: 1503–1510.
- Ritland, K., 1996 A marker-based method for inferences about quantitative inheritance in natural populations. *Evolution* 50: 1062–1073.
- Ritland, K., 2002 Extensions of models for the estimation of mating systems using n independent loci. *Heredity* 88: 221–228.
- Saatchi, M., M. C. McClure, S. D. McKay, M. M. Rolf, J. W. Kim *et al.*, 2011 Accuracies of genomic breeding values in American Angus beef cattle using k-means clustering for cross-validation. *Genet. Sel. Evol.* 43: 40.
- Shen, X., L. Rönnegård, and Ö. Carlborg, 2011 Hierarchical likelihood opens a new way of estimating genetic values using genome-wide dense marker maps. *BMC Proc.* 5(Suppl. 3): S14.
- Shepherd, R., T. Meuwissen, and J. Woolliams, 2010 Genomic selection and complex trait prediction using a fast EM algorithm applied to genome-wide markers. *BMC Bioinformatics* 11: 529.
- Solberg, T. R., A. K. Sonesson, J. A. Woolliams, and T. H. E. Meuwissen, 2009 Reducing dimensionality for prediction of genome-wide breeding values. *Genet. Sel. Evol.* 41: 29.
- Soller, M., 1978 The use of loci associated with quantitative effects in dairy cattle improvement. *Anim. Prod.* 27: 133–139.
- Soller, M., and J. Plotkin-Hazan, 1977 The use of marker alleles for the introgression of linked quantitative alleles. *Theor. Appl. Genet.* 51: 133–137.
- Strandén, I., and O. F. Christensen, 2011 Allele coding in genomic evaluation. *Genet. Sel. Evol.* 43: 25.
- Tibshirani, R., 1996 Regression shrinkage and selection via the lasso. *J. R. Stat. Soc., B* 58: 267–288.
- Usai, M. G., M. E. Goddard, and B. J. Hayes, 2009 LASSO with cross-validation for genomic selection. *Genet. Res. Camb* 91: 427–436.
- VanRaden, P., 2007 Genomic measures of relationship and inbreeding. *Interbull Bull.* 37: 33–36.
- VanRaden, P., 2008 Efficient methods to compute genomic predictions. *J. Dairy Sci.* 91: 4414–4423.
- VanRaden, P. M., and P. G. Sullivan, 2010 International genomic evaluation methods for dairy cattle. *Genet. Sel. Evol.* 42: 7.
- VanRaden, P. M., and G. R. Wiggans, 1991 Derivation, calculation, and use of national animal model information. *J. Dairy Sci.* 74: 2737–2746.
- VanRaden, P. M., C. P. Van Tassell, G. R. Wiggans, T. S. Sonstegard, R. D. Schnabel *et al.*, 2009 Invited review: reliability of genomic predictions for North American Holstein bulls. *J. Dairy Sci.* 92: 16–24.
- Vazquez, A. I., G. J. M. Rosa, K. A. Weigel, G. de los Campos, D. Gianola *et al.*, 2010 Predictive ability of subsets of single nucleotide polymorphisms with and without parent average in US Holsteins. *J. Dairy Sci.* 93: 5942–5949.
- Verbyla, K. L., B. Hayes, P. J. Bowman, and M. E. Goddard, 2009 Short note: accuracy of genomic selection using stochastic search variable selection in Australian Holstein Friesian dairy cattle. *Genet. Res.* 91: 307–311.
- Villanueva, B., R. Pong-Wong, J. Fernández, and M. Toro, 2005 Benefits from marker-assisted selection under an additive polygenic genetic model. *J. Anim. Sci.* 83: 1747–1752.
- Weigel, K. A., G. de los Campos, O. González-Recio, H. Naya, X. L. Wu *et al.*, 2009 Predictive ability of direct genomic values for lifetime net merit of Holstein sires using selected subsets of single nucleotide polymorphism markers. *J. Dairy Sci.* 92: 5248–5257.
- Whittaker, J. C., R. Thompson, and M. C. Denham, 2000 Marker-assisted selection using ridge regression. *Genet. Res.* 75: 249–252.
- Wiggans, G. R., P. M. VanRaden, and T. A. Cooper, 2011 The genomic evaluation system in the United States: past, present, future. *J. Dairy Sci.* 94: 3202–3211.
- Winter, A., W. Krämer, F. A. O. Werner, S. Kollers, S. Kata *et al.*, 2002 Association of a lysine-232/alanine polymorphism in a bovine gene encoding acyl-CoA:diacylglycerol acyltransferase (DGAT1) with variation at a quantitative trait locus for milk fat content. *Proc. Natl. Acad. Sci. USA* 99: 9300–9305.

- Wittenburg, D., N. Melzer, and N. Reinsch, 2011 Including non-additive genetic effects in Bayesian methods for the prediction of genetic values based on genome-wide markers. *BMC Genet.* 12: 74.
- Wolc, A., J. Arango, P. Settar, J. E. Fulton, N. P. O'Sullivan *et al.*, 2011a Persistence of accuracy of genomic estimated breeding values over generations in layer chickens. *Genet. Sel. Evol.* 43: 23.
- Wolc, A., C. Stricker, J. Arango, P. Settar, J. E. Fulton *et al.*, 2011b Breeding value prediction for production traits in layer chickens using pedigree or genomic relationships in a reduced animal model. *Genet. Sel. Evol.* 43: 5.
- Yang, J., B. Benyamin, B. P. McEvoy, S. Gordon, A. K. Henders *et al.*, 2010 Common SNPs explain a large proportion of the heritability for human height. *Nat. Genet.* 42: 565–569.
- Yang, W., and R. J. Tempelman, 2012 A Bayesian antedependence model for whole genome prediction. *Genetics* 190: 1491–1501.
- Yi, N., and S. Banerjee, 2009 Hierarchical generalized linear models for multiple quantitative trait locus mapping. *Genetics* 181: 1101–1113.
- Yu, X., and T. H. E. Meuwissen, 2011 Using the Pareto principle in genome-wide breeding value estimation. *Genet. Sel. Evol.* 43: 35.
- Zhang, Z., J. Liu, X. Ding, P. Bijma, D. J. De Koning *et al.*, 2010 Best linear unbiased prediction of genomic breeding values using a trait-specific marker-derived relationship matrix. *PLoS ONE* 5: e12648.
- Zhao, Y., M. Gowda, F. Longin, T. Würschum, N. Ranc *et al.*, 2012 Impact of selective genotyping in the training population on accuracy and bias of genomic selection. *Theor. Appl. Genet.* DOI: 10.1007/s00122-012-1862-2.
- Zhong, S., J. C. M. Dekkers, R. L. Fernando, and J.-L. Jannink, 2009 Factors affecting accuracy from genomic selection in populations derived from multiple inbred lines: a barley case study. *Genetics* 182: 355–364.
- Zou, F., H. Huang, S. Lee, and I. Hoeschele, 2010 Nonparametric Bayesian variable selection with applications to multiple quantitative trait loci mapping with epistasis and gene–environment interaction. *Genetics* 186: 385–394.
- Zou, H., and T. Hastie, 2005 Regularization and variable selection via the elastic net. *J. R. Stat. Soc. Ser. B Stat. Methodol.* 67: 301–320.

Edited by D. J. de Koning and Lauren M. McIntyre

GENETICS

Supporting Information

<http://www.genetics.org/content/suppl/2012/06/28/genetics.112.143313.DC1>

Whole-Genome Regression and Prediction Methods Applied to Plant and Animal Breeding

Gustavo de los Campos, John M. Hickey, Ricardo Pong-Wong, Hans D. Daetwyler, and Mario P. L. Calus

File S1

Supporting Methods

1. Equivalence between penalized and Bayesian regressions

We show the equivalence first for the **Ridge Regression** (RR) and then for LASSO. The same steps can be used to derive the Bayesian equivalents of other methods, such as Bridge regression. The solution to the optimization problem of the RR is

$$(\hat{\mu}, \hat{\beta}) = \underset{\arg \min}{\left\{ \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 + \lambda \sum_{j=1}^p \beta_j^2 \right\}}.$$

Multiplying the objective function by $-1/2$ and switching from minimization to maximization preserves the solution, therefore:

$$(\hat{\mu}, \hat{\beta}) = \underset{\arg \max}{\left\{ -\frac{1}{2} \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 - \frac{1}{2} \lambda \sum_{j=1}^p \beta_j^2 \right\}}.$$

Further, dividing the objective function by any positive constant preserves the solution;

therefore for any $\sigma^2 > 0$ we have,

$$(\hat{\mu}, \hat{\beta}) = \underset{\arg \max}{\left\{ -\frac{1}{2\sigma^2} \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 - \frac{\lambda}{2\sigma^2} \sum_{j=1}^p \beta_j^2 \right\}}.$$

Moreover, for any positive value of σ_β^2 such that $\lambda = \sigma^2 \sigma_\beta^{-2}$ we have:

$$(\hat{\mu}, \hat{\beta}) = \underset{\arg \max}{\left\{ -\frac{1}{2\sigma^2} \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 - \frac{1}{2\sigma_\beta^2} \sum_{j=1}^p \beta_j^2 \right\}}.$$

Finally, applying any monotonic transformation to the objective function also preserves the solution, therefore:

$$\begin{aligned}
(\hat{\mu}, \hat{\boldsymbol{\beta}}) &= \underset{\arg \max}{\left\{ -\frac{1}{2\sigma^2} \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 - \frac{1}{2\sigma_\beta^2} \sum_{j=1}^p \beta_j^2 \right\}} \\
&= \underset{\arg \max}{\left\{ \exp \left[-\frac{1}{2\sigma^2} \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 - \frac{1}{2\sigma_\beta^2} \sum_{j=1}^p \beta_j^2 \right] \right\}} \\
&= \underset{\arg \max}{\left\{ \exp \left[-\frac{1}{2\sigma^2} \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 \right] \exp \left[-\frac{1}{2\sigma_\beta^2} \sum_{j=1}^p \beta_j^2 \right] \right\}} \\
&= \underset{\arg \max}{\left\{ \prod_{i=1}^n \exp \left(-\frac{\left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2}{2\sigma^2} \right) \right\} \left\{ \prod_{j=1}^p \exp \left(-\frac{\beta_j^2}{2\sigma_\beta^2} \right) \right\}}
\end{aligned}$$

The first term in the above expression, is proportional to a Gaussian likelihood for data (y_i) with mean $\mu + \sum_{j=1}^p x_{ij} \beta_j$ and residual variance σ^2 . And the second term is proportional to a Gaussian prior for marker effects with mean equal to zero and variance σ_β^2 . Specifically, the solution to RR optimization problem is equivalent to the posterior mode of the following Bayesian model:

$$\begin{aligned}
p(\mathbf{y}, \boldsymbol{\beta}, \mu | \sigma^2, \sigma_\beta^2) &\propto p(\mathbf{y} | \boldsymbol{\beta}, \mu, \sigma^2, \sigma_\beta^2) p(\boldsymbol{\beta} | \sigma_\beta^2) \\
&\propto \prod_{i=1}^n N(y_i | \mu + \sum_{j=1}^p x_{ij} \beta_j, \sigma^2) \prod_{j=1}^p N(\beta_j^2 | 0, \sigma_\beta^2).
\end{aligned}$$

A similar reasoning can be used to show the equivalence for the LASSO and in general for Bridge regression. For the LASSO, we replace the penalty $\sum_{j=1}^p \beta_j^2$ with $\sum_{j=1}^p |\beta_j|$; therefore

$$\begin{aligned}
(\hat{\mu}, \hat{\beta})_{\arg \min} &= \left\{ \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 + \lambda \sum_{j=1}^p |\beta_j| \right\} \\
&= \left\{ -\frac{1}{2} \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 - \frac{\lambda}{2} \sum_{j=1}^p |\beta_j| \right\} \\
&= \left\{ -\frac{1}{2\sigma^2} \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 - \frac{\lambda}{2\sigma^2} \sum_{j=1}^p |\beta_j| \right\}, \text{ for any } \sigma^2 > 0 \\
&= \left\{ -\frac{1}{2\sigma^2} \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 - \tilde{\lambda} \sum_{j=1}^p |\beta_j| \right\}, \text{ for any } \tilde{\lambda} > 0, \text{ such that } \tilde{\lambda} = \frac{\lambda}{2\sigma^2} \\
&= \left\{ \exp \left\{ -\frac{1}{2\sigma^2} \sum_i \left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2 \right\} \exp \left\{ -\tilde{\lambda} \sum_{j=1}^p |\beta_j| \right\} \right\} \\
&= \left\{ \prod_{i=1}^n \exp \left\{ -\frac{\left(y_i - \mu - \sum_{j=1}^p x_{ij} \beta_j \right)^2}{2\sigma^2} \right\} \prod_{j=1}^p \exp \left\{ -\tilde{\lambda} |\beta_j| \right\} \right\}
\end{aligned}$$

As with the RR, the first term is proportional to the following Gaussian likelihood. The second term is proportional to the product of p IID Double-Exponential, or Laplace, priors densities for marker effects.

2. On the joint density of genetic values of genotyped and un-genotyped individuals

In this section we consider the problem of deriving the joint density of genetic values when some individuals (set 1) were not genotyped and others (set 2) were genotyped. We show that the joint density of the genetic values of these two sets of individuals, denoted as \mathbf{g}_1 and \mathbf{g}_2 , respectively, in the RR-BLUP model is a mixture of multivariate normal densities.

When all individuals are genotyped. Following standard assumptions, the marginal distribution of genomic values in RR-BLUP is as follows:

$$p\left(\begin{matrix} \mathbf{g}_1 \\ \mathbf{g}_2 \end{matrix} \middle| \mathbf{X}_1, \mathbf{X}_2, \sigma_u^2\right) = MVN\left(\begin{matrix} \mathbf{0} \\ \mathbf{0} \end{matrix}, \begin{bmatrix} \mathbf{X}_1\mathbf{X}_1' & \mathbf{X}_1\mathbf{X}_2' \\ \mathbf{X}_2\mathbf{X}_1' & \mathbf{X}_2\mathbf{X}_2' \end{bmatrix} \sigma_u^2\right), \quad [1]$$

where, \mathbf{X}_1 and \mathbf{X}_2 are matrices of marker genotypes and σ_u^2 is a variance parameter and MVN denotes a multivariate normal density.

When some individuals are not genotyped. Consider the case where only individuals in set 2 are genotyped. In this case, we need to derive the joint density of genetic values given \mathbf{X}_2 , pedigree relationships (denoted as P) and σ_u^2 that is, $p(\mathbf{g}_1, \mathbf{g}_2 | \mathbf{X}_2, P, \sigma_u^2)$. To derive this density we first augment the probability model by introducing \mathbf{X}_1 , and subsequently integrating it out:

$$\begin{aligned} p(\mathbf{g}_1, \mathbf{g}_2 | \mathbf{X}_2, P, \sigma_u^2) &= \int p(\mathbf{g}_1, \mathbf{g}_2, \mathbf{X}_1 | \mathbf{X}_2, P, \sigma_u^2) d\mathbf{X}_1 \\ &= \int p(\mathbf{g}_1, \mathbf{g}_2 | \mathbf{X}_1, \mathbf{X}_2, \sigma_u^2) p(\mathbf{X}_1 | \mathbf{X}_2, P) d\mathbf{X}_1 \\ &= \int MVN\left(\begin{matrix} \mathbf{g}_1 \\ \mathbf{g}_2 \end{matrix} \middle| \begin{matrix} \mathbf{0} \\ \mathbf{0} \end{matrix}, \begin{bmatrix} \mathbf{X}_1\mathbf{X}_1' & \mathbf{X}_1\mathbf{X}_2' \\ \mathbf{X}_2\mathbf{X}_1' & \mathbf{X}_2\mathbf{X}_2' \end{bmatrix} \sigma_u^2\right) p(\mathbf{X}_1 | \mathbf{X}_2, P) d\mathbf{X}_1 \end{aligned} \quad [2]$$

The first density on the right-hand side is simply the MVN density of expression [1]. The second component of the right-hand side, $p(\mathbf{X}_1 | \mathbf{X}_2, P)$ gives the probability density function of the unknown-genotypes given the observed genotypes and the pedigree. This is the density we would use, for instance, in pedigree-based imputation algorithms. For every realization of \mathbf{X}_1 we have a peculiar MVN with a particular co-variance structure (see right-hand side of expression [2]). Therefore, we conclude that the joint density of genetic values is a mixture of scaled-multivariate normal densities.

Existing proposals for joint analysis of genotyped and un-genotyped individuals (e.g., CHRISTENSEN and LUND 2010; AGUILAR *et al.* 2010) assume that the joint density of genetic values of these two groups of individuals is MVN. In light of the above-results, these methods should be considered linear approximation to a non-linear problem.